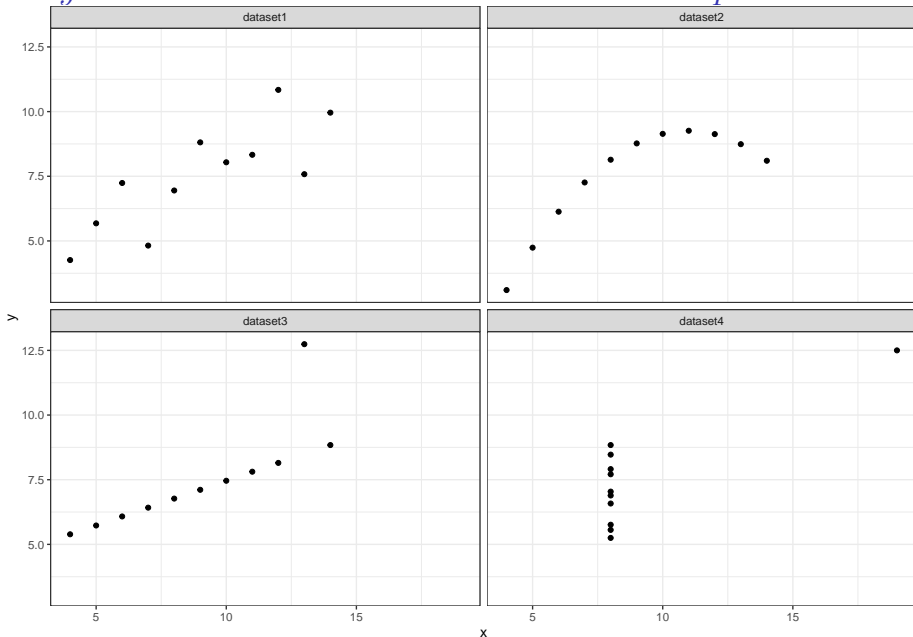


R Data Vis: ggplot

Brian Wright, PhD

Feb, 2020

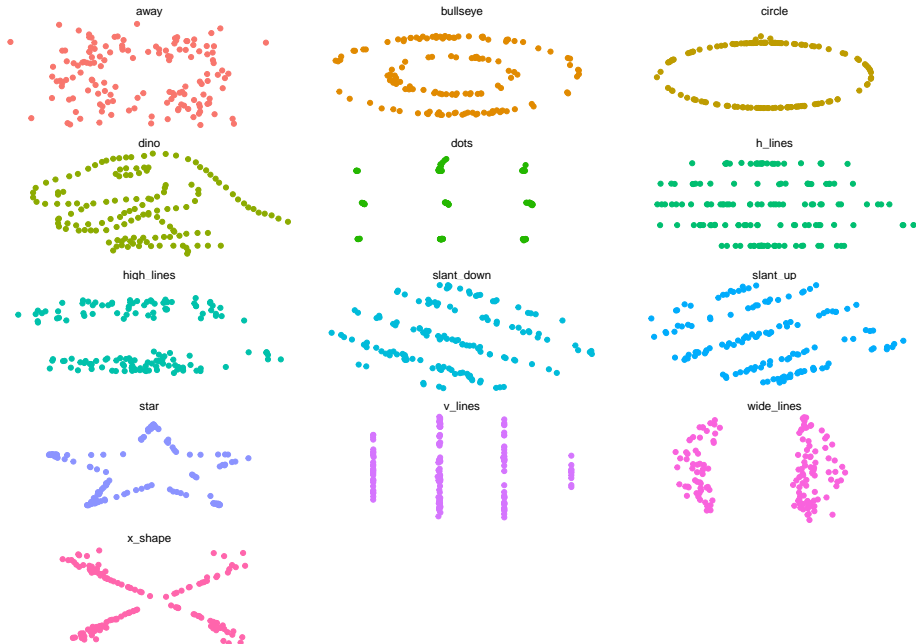
Why visualize data? Anscombe's data example



All Have the Summary Stats

Statistic	Value
mean(x)	9
mean(y)	7.5
var(x)	11
var(y)	4.13
cor(x,y)	0.82

The DataSaurus dozen



Same Stats

Statistic	Value
mean(x)	54.3
mean(y)	47.8
var(x)	281
var(y)	725
cor(x,y)	-0.07

Bottom line

- Summary statistics cannot always distinguish datasets
- Take advantage of humans' ability to visually recognize and remember patterns
- Find discrepancies in the data more easily

What is ggplot2?

- A second (and final) iteration of the ggplot
- Implementation of Wilkerson's Grammar of Graphics in R
- Conceptually, a way to layer different elements onto a canvas to create a data visualization
- Started as Dr. Hadley Wickham's PhD thesis (with Dr. Dianne Cook)
- Won the John M. Chambers Statistical Software Award in 2006
- Mimicked in other software platforms
 - ▶ ggplot and seaborn in Python
 - ▶ Translated in plotly

*ggplot2 uses the **grammar of graphics***

A grammar ...

- compose and re-use small parts
- build complex structures from simpler units

of graphics

- think of yourself as a painter
- build a visualization using layers on a canvas
- draw layers on top of each other

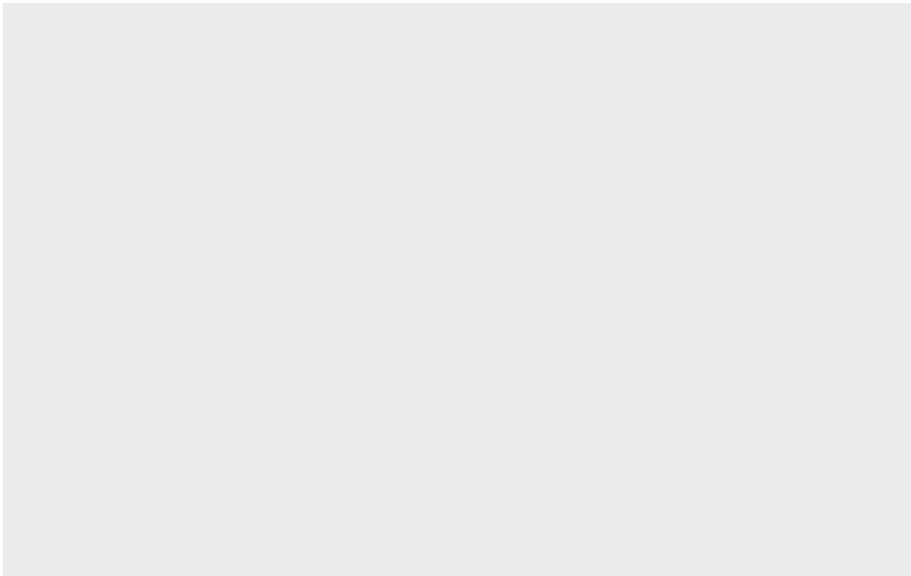
A dataset

```
library(tidyverse)
library(rio)
beaches <- import('beaches.csv')
```

```
#> # A tibble: 6 x 12
#>   date       year month   day season rainfall temperature enterococci day_num
#>   <date>     <int> <int> <int> <int>    <dbl>      <dbl>      <dbl>    <int>
#> 1 2013-01-02  2013     1     2     1      0        23.4        6.7      2
#> 2 2013-01-06  2013     1     6     1      0        30.3         2      6
#> 3 2013-01-12  2013     1    12     1      0        31.4       69.1    12
#> 4 2013-01-18  2013     1    18     1      0        46.4         9    18
#> 5 2013-01-24  2013     1    24     1      0        27.5       33.9    24
#> 6 2013-01-30  2013     1    30     1      0.6       26.6       26.5    30
#> # ... with 3 more variables: month_num <int>, month_name <chr>,
#> #   season_name <chr>
```

Building a graph: Start with a blank canvas

```
ggplot()
```



Add a data set

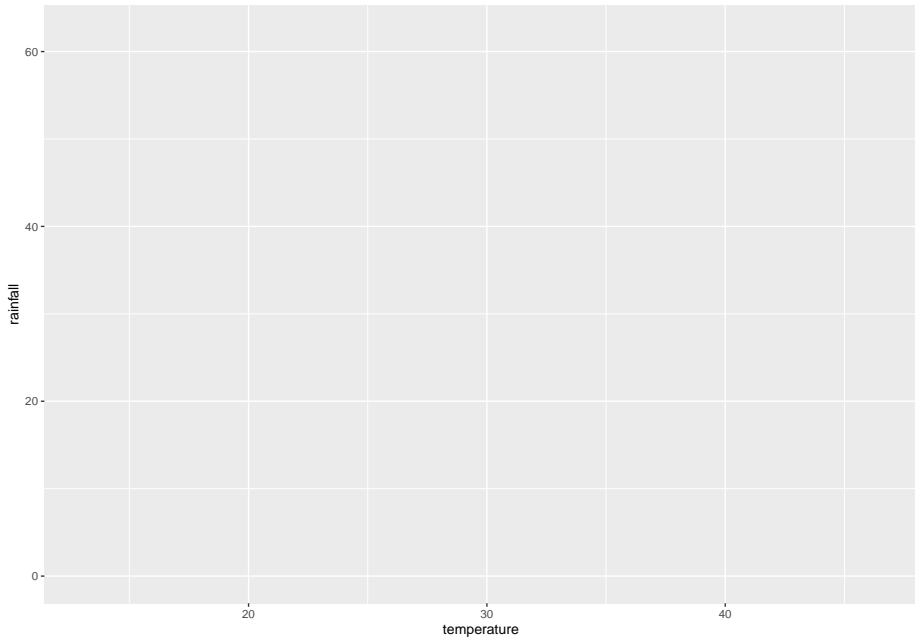
```
ggplot(  
  data = beaches #<< loaded earlier  
)
```

Add a mapping from data to elements

```
ggplot(  
  data = beaches,  
  mapping = aes( # everytime we add a "data element"  
    #we add a aesthetic  
    x = temperature,  
    y = rainfall  
  )  
)
```

What goes in

- the x and y axes
- the color of markers
- the shape of markers

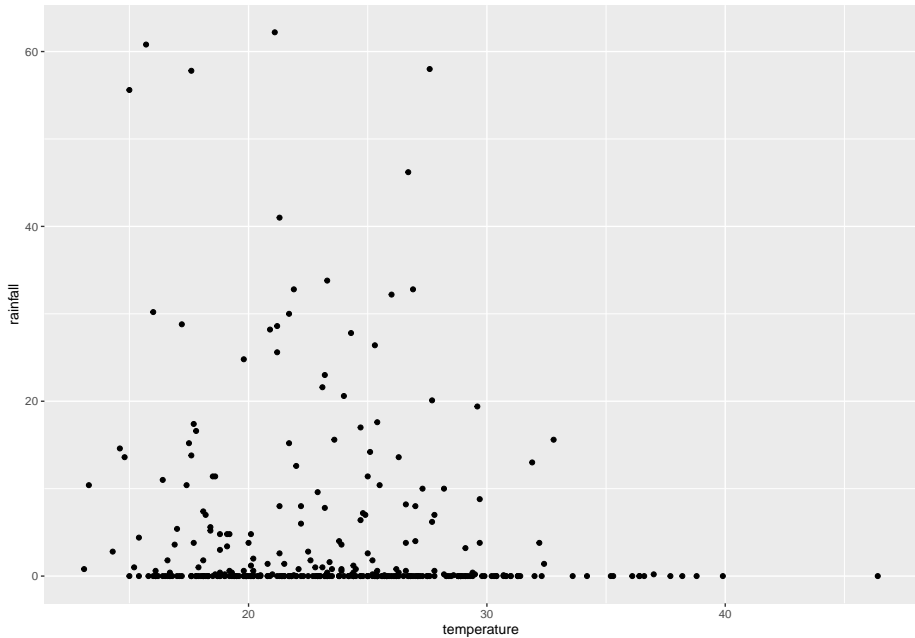


Add a geometry to draw

```
ggplot(  
  data = beaches,  
  mapping = aes(  
    x = temperature,  
    y = rainfall  
  )  
) +  
  geom_point() #?
```

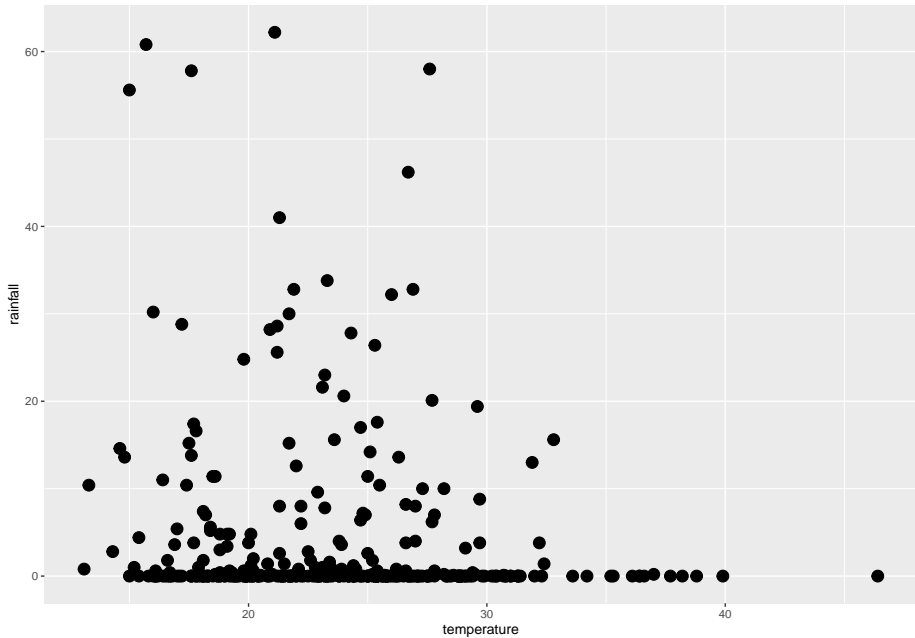
What to draw:

- Points, lines
- Histogram, bars, pies, etc.



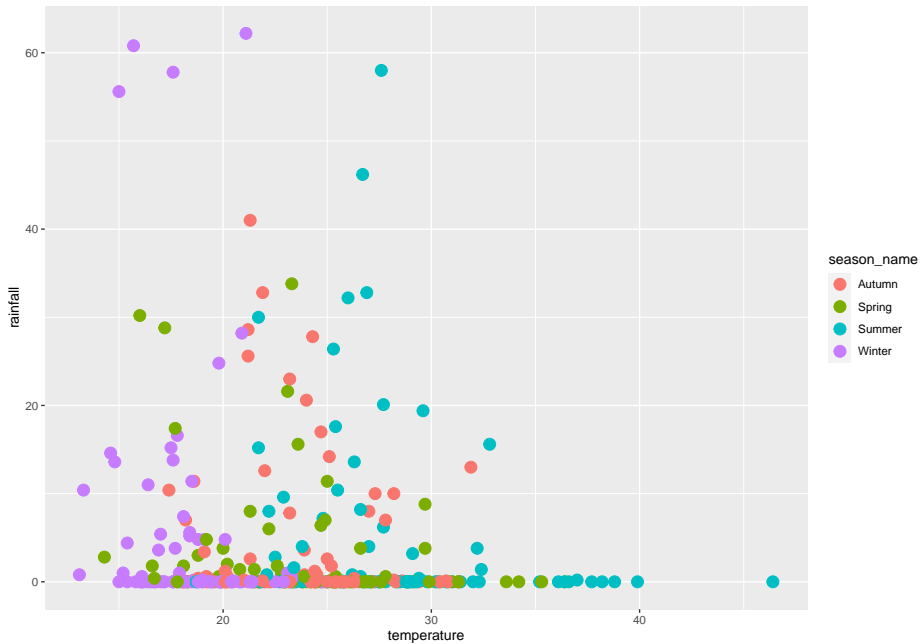
Add options for the geom

```
ggplot(  
  data = beaches,  
  mapping = aes(  
    x = temperature,  
    y = rainfall  
  )  
) +  
  geom_point(size = 4)#?
```

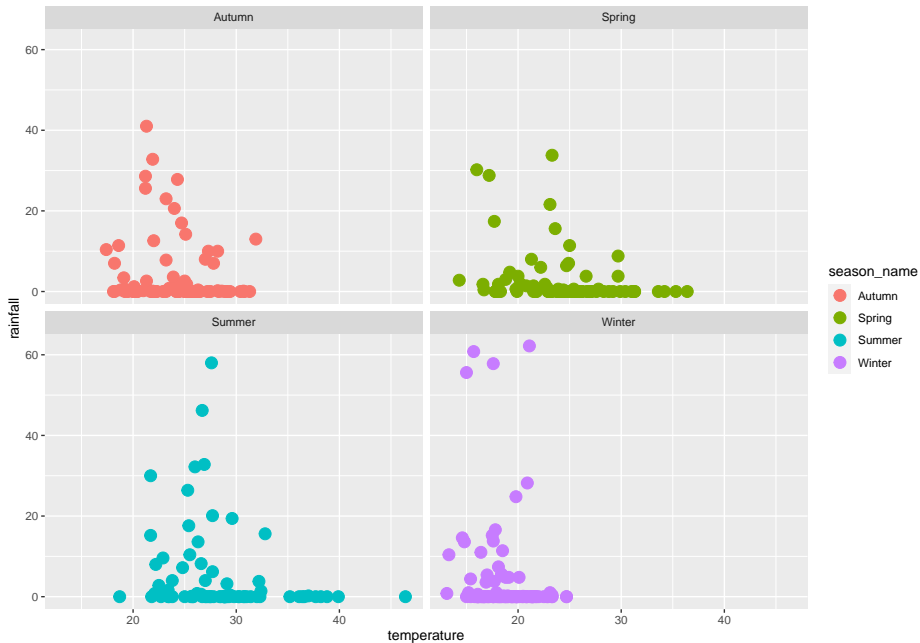
Add a mapping to modify the geom

```
ggplot(  
  data = beaches,  
  mapping = aes(  
    x = temperature,  
    y = rainfall  
  )  
) +  
  geom_point(  
    mapping = aes(color = season_name),  
    size = 4 #Why do we need the mapping?  
  )
```



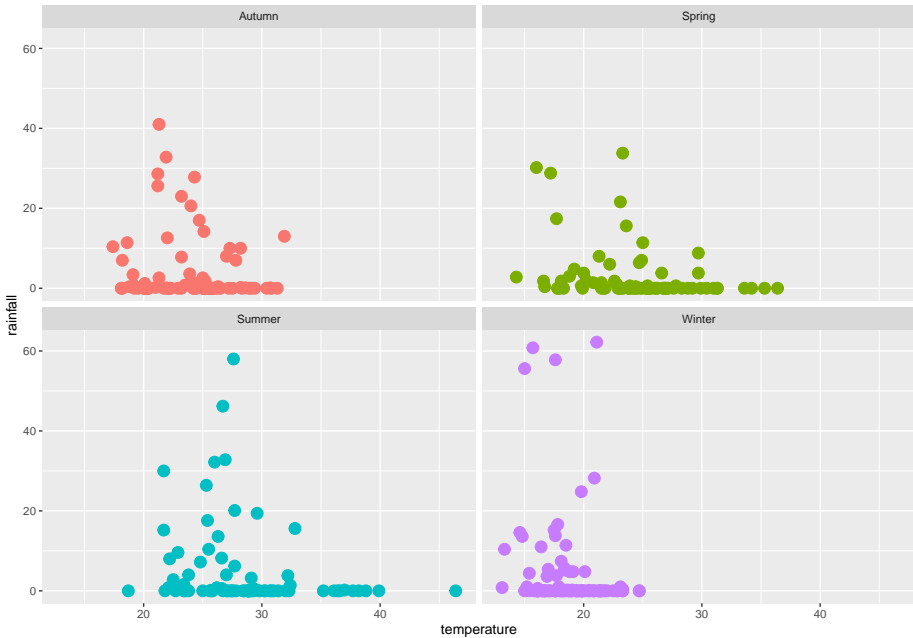
Split into facets

```
ggplot(  
  data = beaches,  
  mapping = aes(  
    x = temperature,  
    y = rainfall  
  )  
) +  
  geom_point(  
    mapping = aes(color = season_name),  
    size = 4  
  ) +  
  facet_wrap( ~ season_name) ###
```



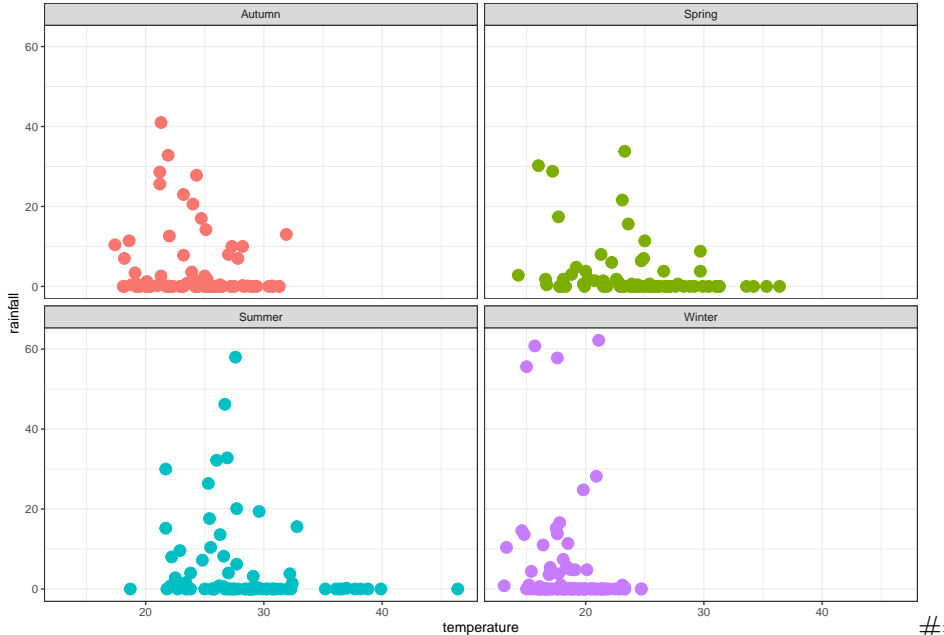
Remove the legend

```
ggplot(  
  data = beaches,  
  mapping = aes(  
    x = temperature,  
    y = rainfall  
  )  
) +  
  geom_point(  
    mapping = aes(color = season_name),  
    size = 4,  
    show.legend = FALSE ###  
  ) +  
  facet_wrap( ~ season_name)
```

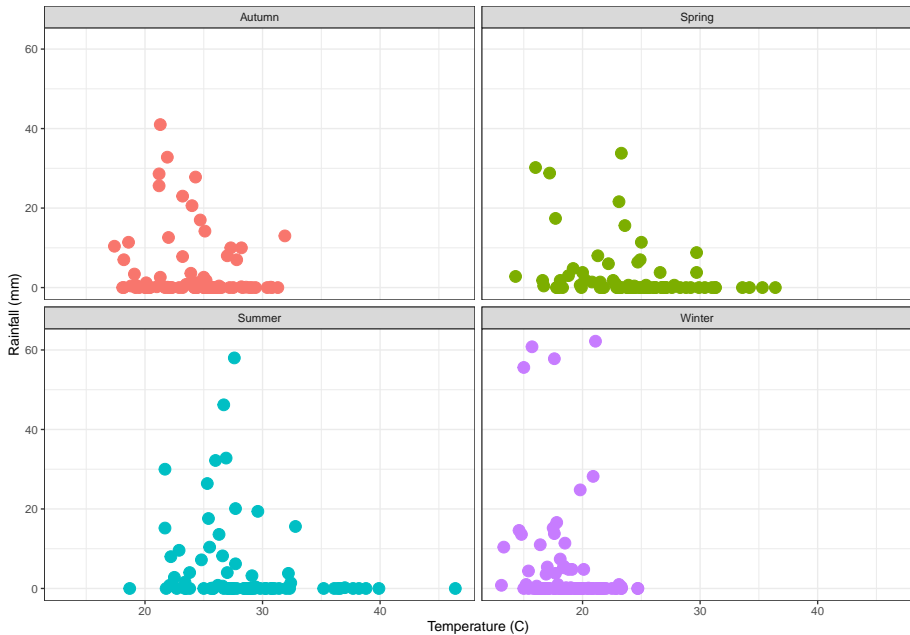


Change the general theme

```
ggplot(  
  data = beaches,  
  mapping = aes(  
    x = temperature,  
    y = rainfall  
  )  
) +  
  geom_point(  
    mapping = aes(color = season_name),  
    size = 4,  
    show.legend = FALSE  
  ) +  
  facet_wrap( ~ season_name) +  
  theme_bw() ###
```

Update the labels

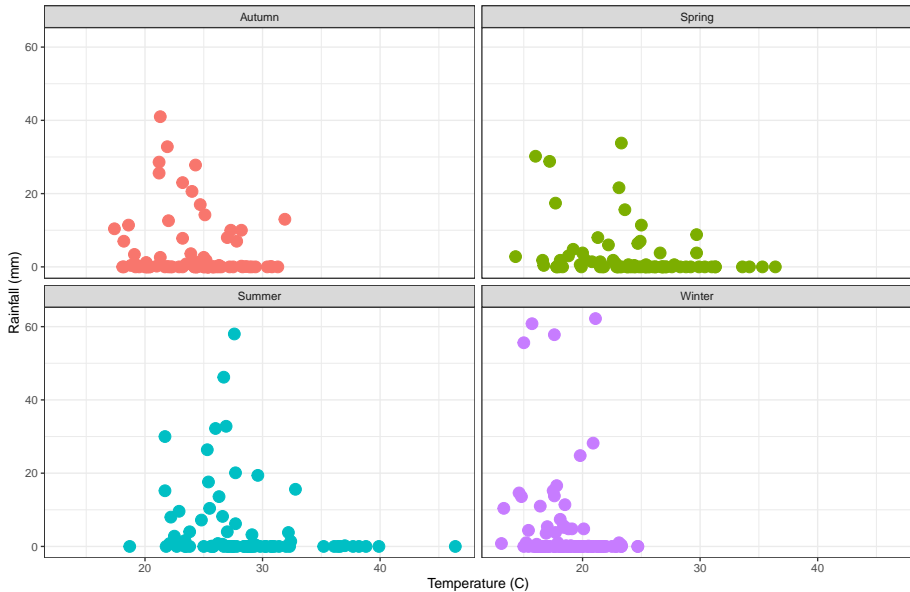


Add titles

```
ggplot(  
  data = beaches,  
  mapping = aes(  
    x = temperature,  
    y = rainfall)  
) +  
  geom_point(  
    mapping = aes(color = season_name),  
    size = 4,  
    show.legend = FALSE  
) +  
  facet_wrap( ~ season_name) +  
  theme_bw() +  
  labs(x = 'Temperature (C)',  
       y = 'Rainfall (mm)',  
       title = 'Sydney weather by season', ###  
       subtitle = "Data from 2013 to 2018") ###
```

Sydney weather by season

Data from 2013 to 2018

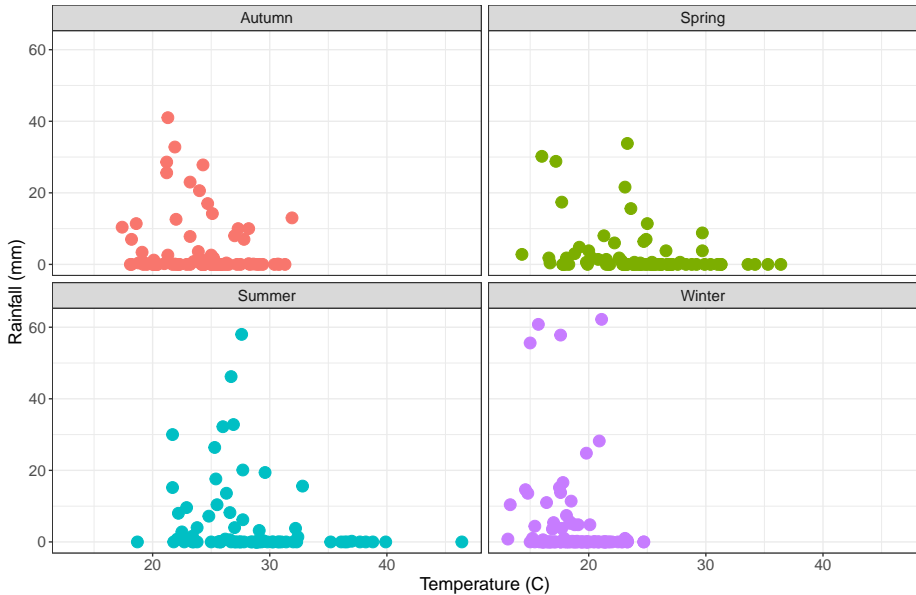


Customize

```
ggplot(  
  data = beaches,  
  mapping = aes(  
    x = temperature,  
    y = rainfall  
  )  
) +  
  geom_point(  
    mapping = aes(color = season_name),  
    size = 4,  
    show.legend = FALSE  
  ) +  
  facet_wrap( ~ season_name) +  
  theme_bw() +  
  labs(x = 'Temperature (C)',  
       y = 'Rainfall (mm)',  
       title = 'Sydney weather by season',  
       subtitle = "Data from 2013 to 2018") +  
  theme(axis.title = element_text(size = 14), ###  
        axis.text = element_text(size = 12), ###  
        strip.text = element_text(size = 12)) ###
```

Sydney weather by season

Data from 2013 to 2018



The grammar

- Data
- Aesthetics (or aesthetic mappings)
- Geometries (as layers) or Statistics (as computed layers)
- Facets
- Themes
- (Coordinates)
- (Scales)

Peeking under the hood ...

We input the below items

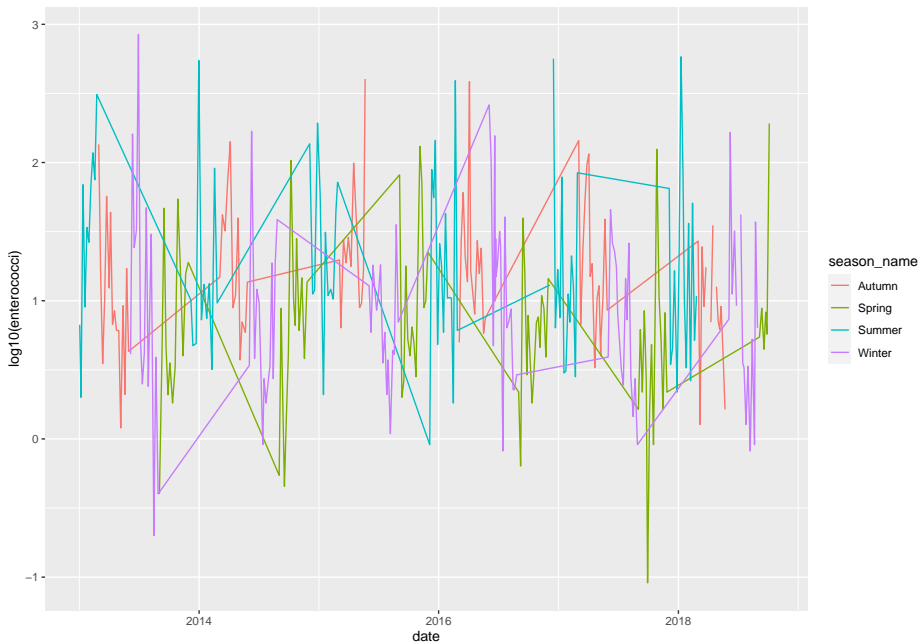
```
ggplot(  
  data = beaches,  
  aes(x = temperature,  
       y = rainfall)  
) +  
  geom_point()
```


What's really run is ...

```
ggplot(  
  data = beaches,  
  mapping = aes(  
    x = temperature, y = rainfall)  
  ) +  
layer(  
  geom = "point",  
  stat = "identity",  
  position = "identity") +  
facet_null() +  
theme_grey() +  
coord_cartesian() +  
scale_x_continuous() +  
scale_y_continuous()
```

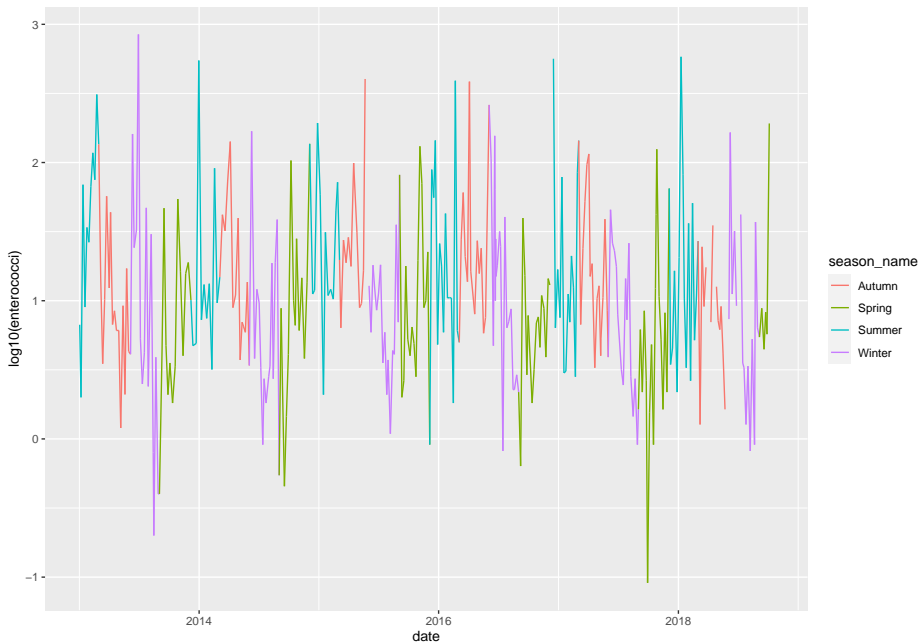
Exploring aesthetics: Mapping color

```
ggplot(  
  data=beaches,  
  aes(x = date,  
      y = log10(enterococci),  
      color=season_name)  
) +  
  geom_line()
```



Adding groups to the mapping

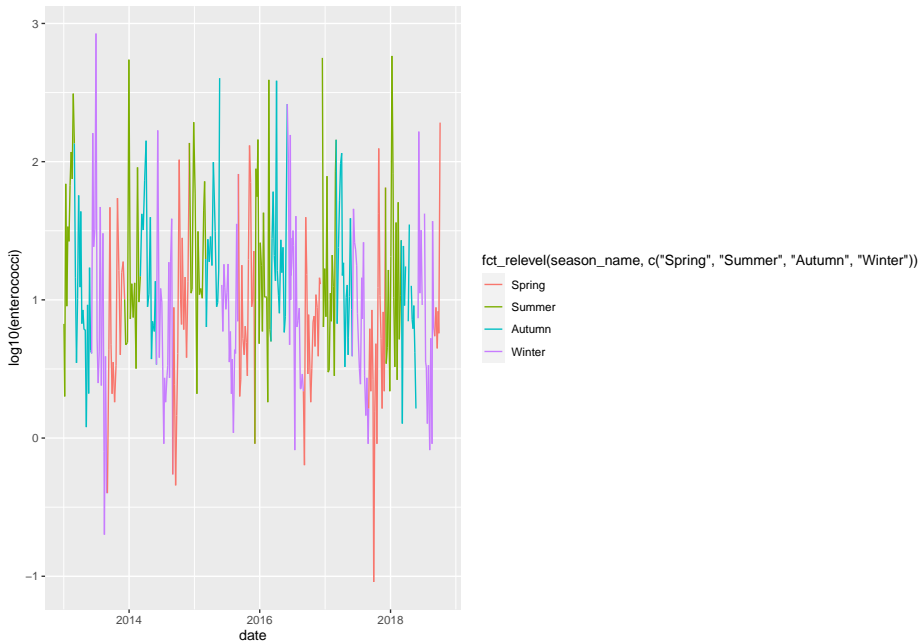
```
ggplot(  
  data=beaches,  
  aes(x = date,  
      y = log10(enterococci),  
      color = season_name,  
      group = 1) ###  
) +  
  geom_line()
```



Fixing the legend ordering

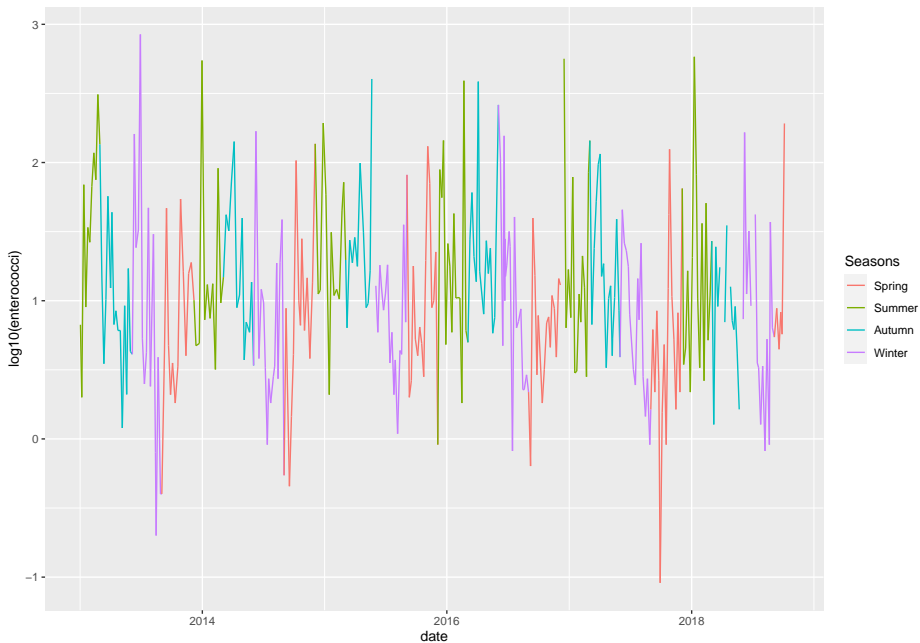
```
ggplot(  
  data=beaches,  
  aes(x = date,  
      y = log10(enterococci),  
      color = fct_relevel(season_name,  
                          c('Spring', 'Summer', 'Autumn', 'Winter')),  
      group = 1)  
) +  
  geom_line()
```

Yikes!



Fixing the legend ordering

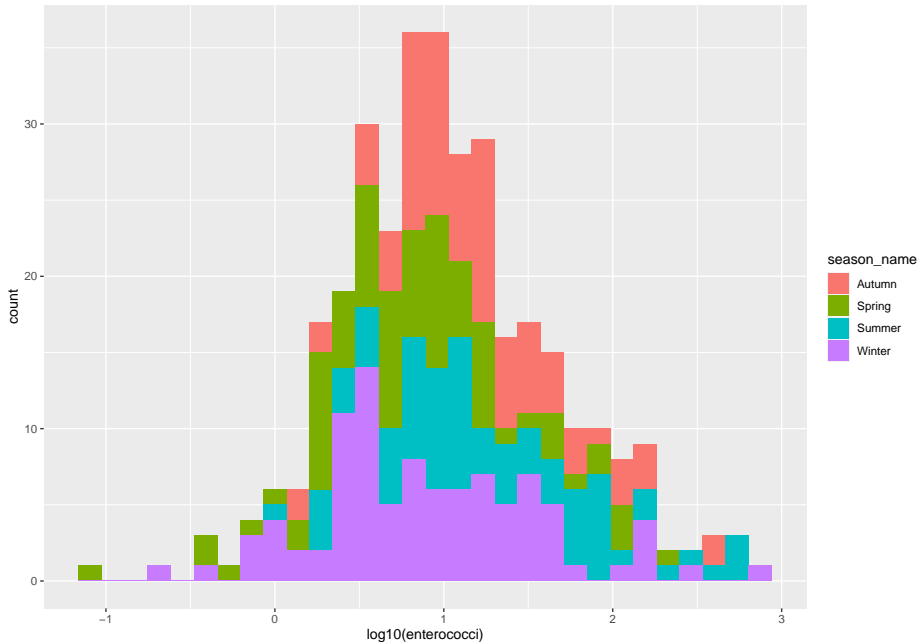
```
ggplot(  
  data=beaches,  
  aes(x = date,  
      y = log10(enterococci),  
      color = fct_relevel(season_name,  
                          c('Spring', 'Summer', 'Autumn', 'Winter')),  
      group = 1)  
) +  
  geom_line() +  
  labs(color = 'Seasons') ###
```

You can also fill based on data

```
ggplot(  
  data=beaches,  
  aes(x = log10(enterococci),  
      fill = season_name)  
)+  
geom_histogram()
```

Works a little better

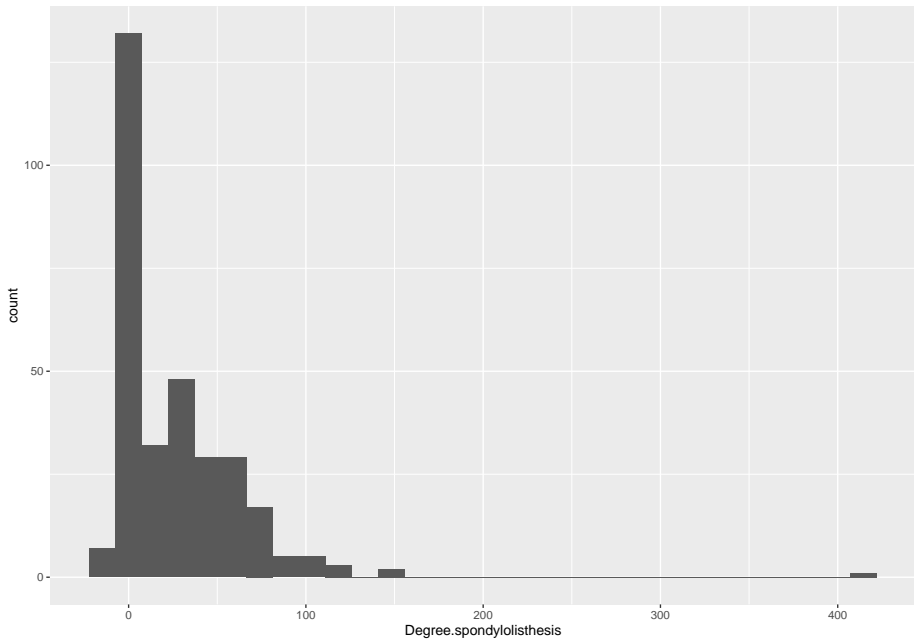


Exploring geometries: Univariate plots

```
library(tidyverse)
library(rio)
dat_spine <- import('Dataset_spine.csv',
                    check.names = T)

ggplot(
  data=dat_spine,
  aes(x = Degree.spondylolisthesis)
)+
  geom_histogram()
```

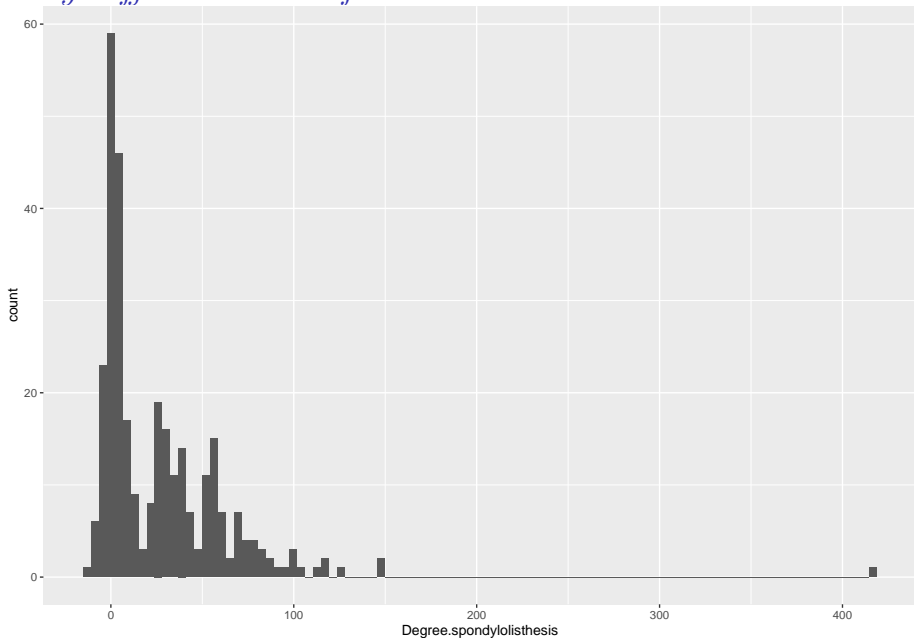
```
#> `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



Histograms

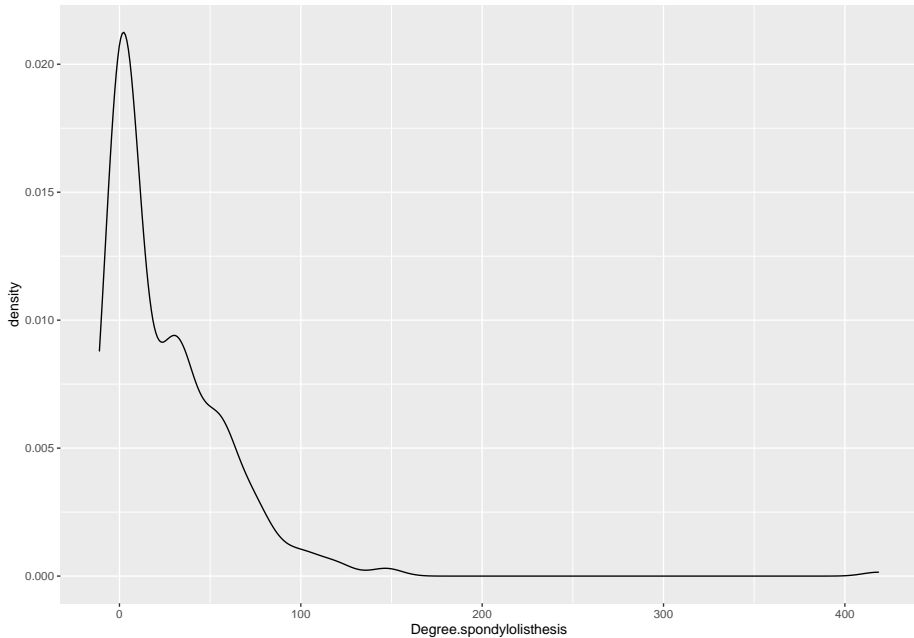
```
ggplot(  
  data=dat_spine,  
  aes(x = Degree.spondylolisthesis)  
) +  
  geom_histogram(bins = 100)
```

Very different view of the data



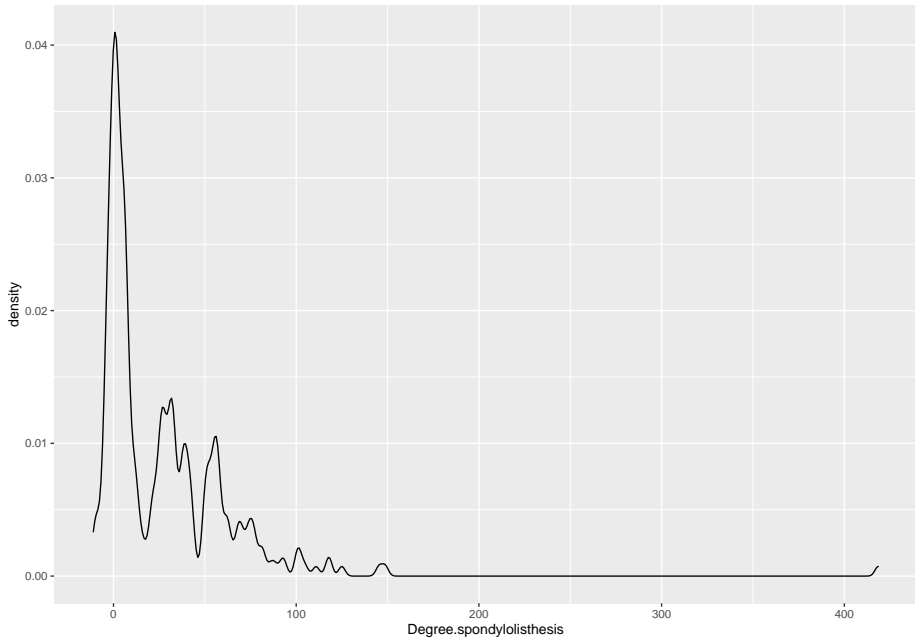
Density plots

```
ggplot(  
  data=dat_spine,  
  aes(x = Degree.spondylolisthesis)  
) +  
  geom_density()
```

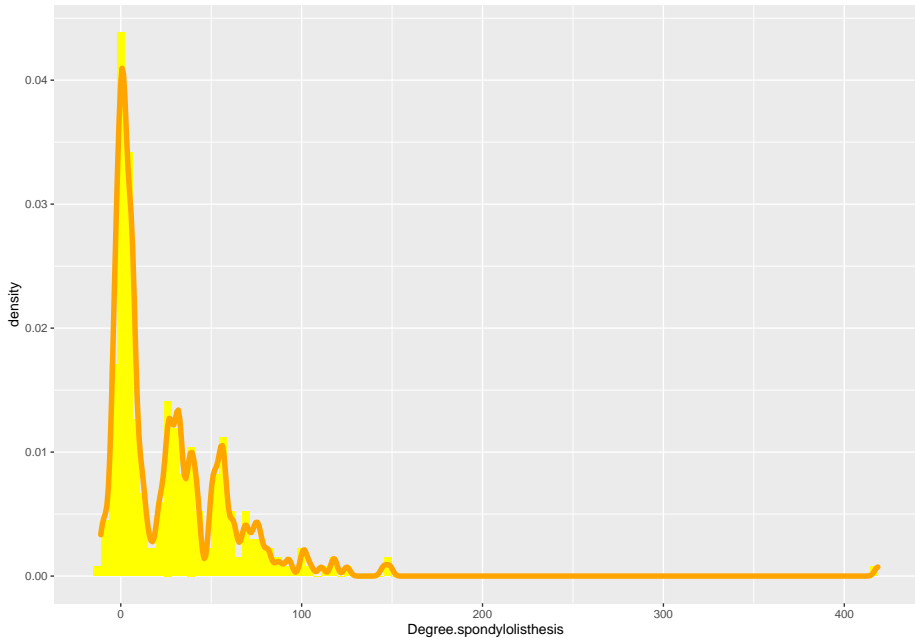
Density plots

```
ggplot(  
  data=dat_spine,  
  aes(x = Degree.spondylolisthesis)  
)+  
geom_density(adjust = 1/5) # Use 1/5 the bandwidth
```



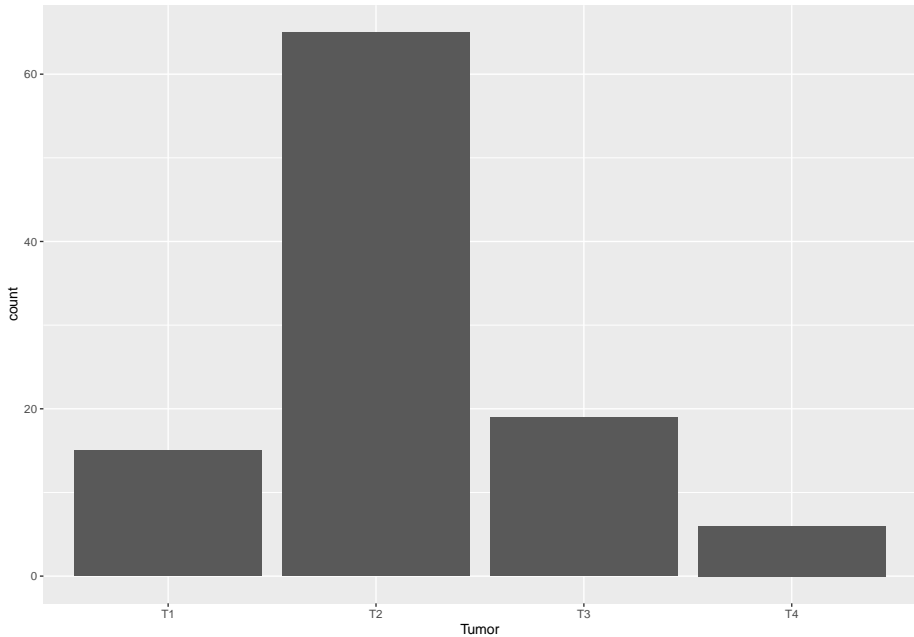
Layering geometries

```
ggplot(  
  data=dat_spine,  
  aes(x = Degree.spondylolisthesis,  
      y = stat(density))  
) + # Re-scales histogram  
  geom_histogram(bins = 100, fill='yellow') +  
  geom_density(adjust = 1/5, color = 'orange', size = 2)
```



Bar plots (categorical variable)

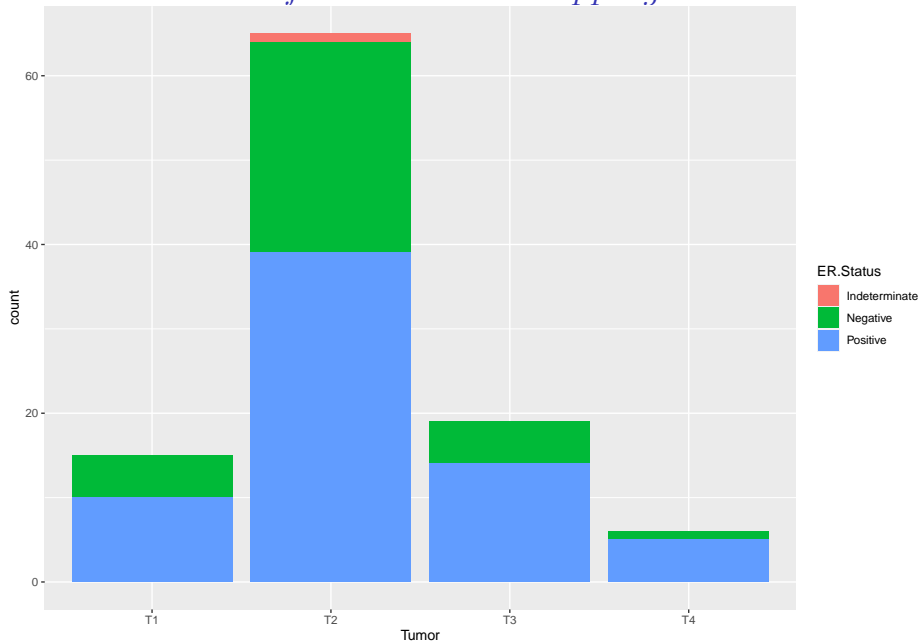
```
dat_brca <-  
  rio::import('clinical_data_breast_cancer_modified.csv',  
              check.names = T)  
  
ggplot(  
  data=dat_brca,  
  aes(x = Tumor)  
)+  
  geom_bar()
```



Bar plots (categorical variable)

```
ggplot(  
  data=dat_brca,  
  aes(x = Tumor,  
       fill = ER.Status)  
)+ #<<  
geom_bar()
```

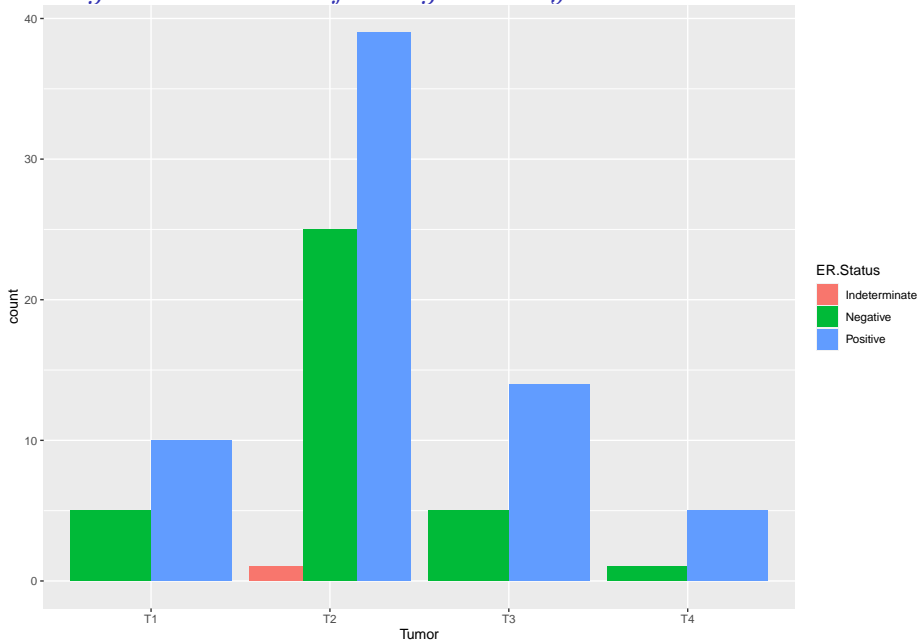

Add additional information via mapping



Bar plots (categorical variable)

```
ggplot(  
  data=dat_brca,  
  aes(x = Tumor,  
       fill = ER.Status)  
) +  
  geom_bar(position = 'dodge')  
  # Default is position = "stack"
```

Change the nature of the geometry



Graphing tabulated data

```
(tabulated <- dat_brca %>% count(Tumor))
```

```
#> # A tibble: 4 x 2
#>   Tumor      n
#>   <chr> <int>
#> 1 T1      15
#> 2 T2      65
#> 3 T3      19
#> 4 T4       6
```

```
ggplot(
  data = tabulated,
  aes(x = Tumor, y = n)
) +
  geom_bar()
```

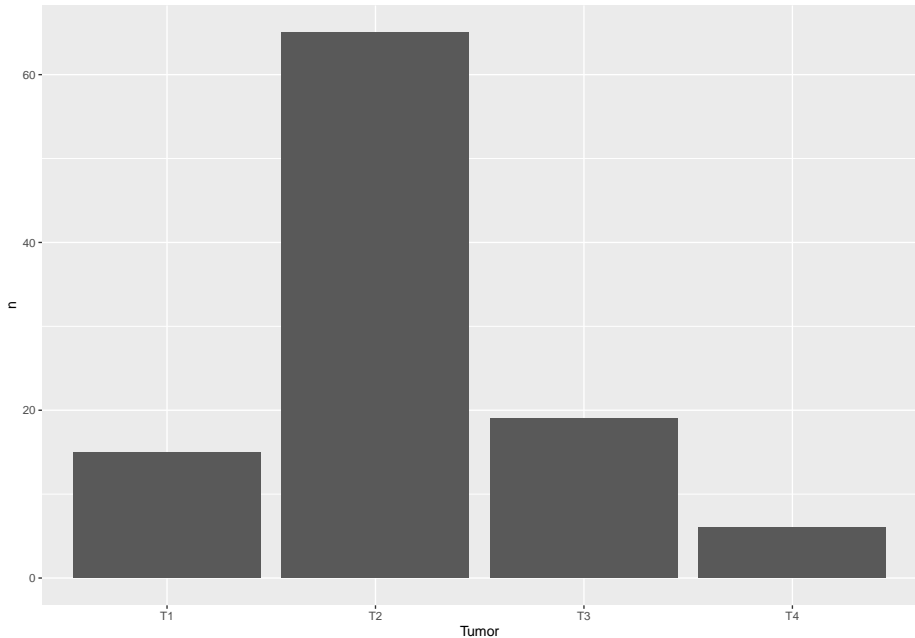
```
#> Error: stat_count() can only have an x or y aesthetic.
```

Graphing tabulated data

```
tabulated <- dat_brca %>% count(Tumor)
tabulated

ggplot(
  data = tabulated,
  aes(x = Tumor, y = n)
) +
geom_bar(stat = 'identity') ###
```

- Here we need to change the default computation
- The barplot usually computes the counts (stat_count)
- We suppress that here since we have already done the computation



Peeking under the hood

```
plt <- ggplot(  
  data = tabulated,  
  aes(x = Tumor, y = n)  
) +  
  geom_bar()
```

```
plt$layers
```

```
#> [[1]]  
#> geom_bar: width = NULL, na.rm = FALSE, orientation = NA  
#> stat_count: width = NULL, na.rm = FALSE, orientation = NA  
#> position_stack
```

Peeking under the hood

```
plt <- ggplot(  
  data = tabulated,  
  aes(x = Tumor, y = n)) +  
  geom_bar(stat = 'identity')
```

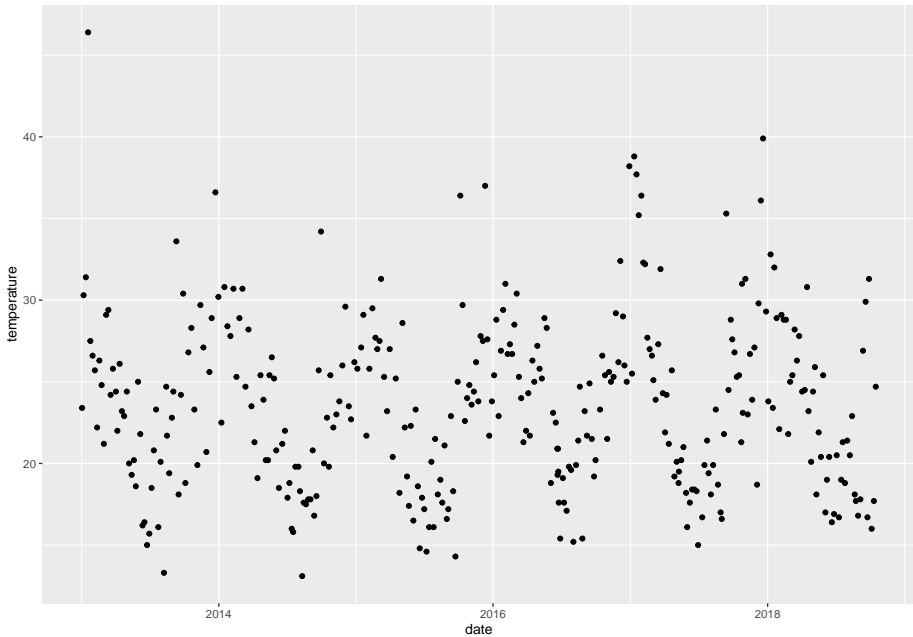
```
plt$layers
```

```
#> [[1]]  
#> geom_bar: width = NULL, na.rm = FALSE, orientation = NA  
#> stat_identity: na.rm = FALSE  
#> position_stack
```

Each layer has a geometry, statistic and position associated with it

Bivariate plots: Scatter plots

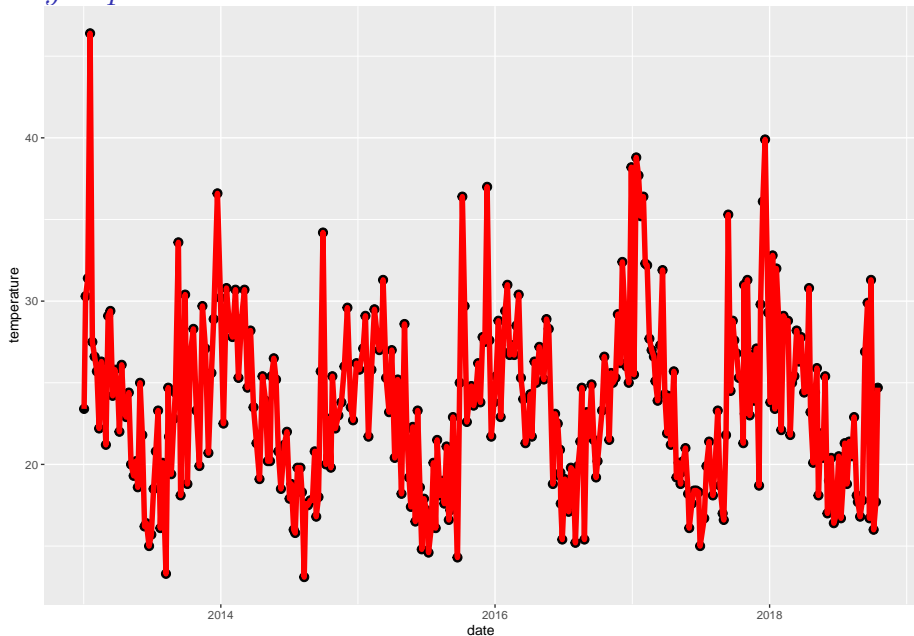
```
ggplot(  
  data = beaches,  
  aes(x = date, y = temperature)  
) +  
  geom_point()
```



Scatter plots

```
ggplot(  
  data = beaches,  
  aes(x = date, y = temperature, group=1) #Add the group argu.  
) +  
  geom_point(color='black', size = 3) +  
  geom_line(color='red', size=2)
```

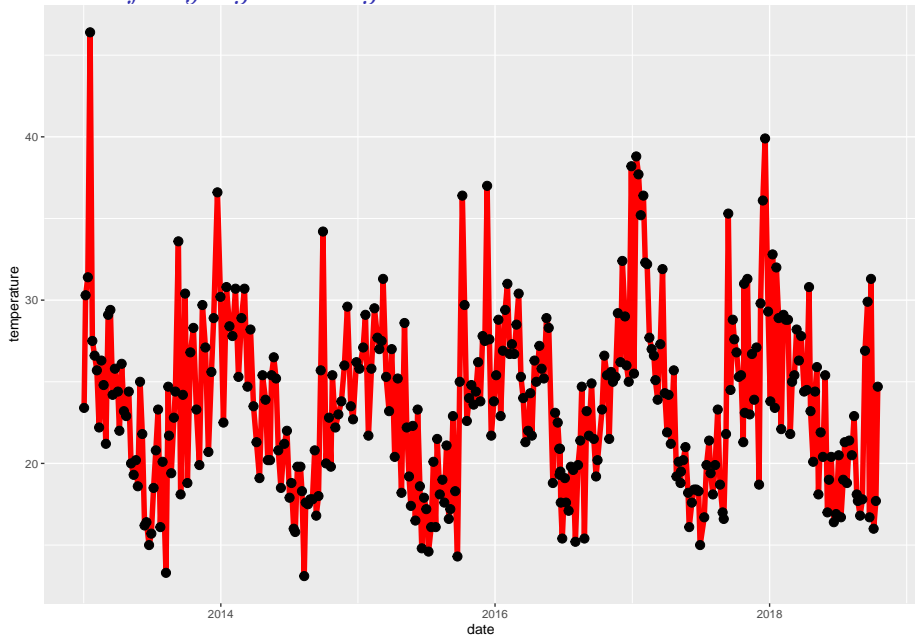
Layer points and lines



Scatter plots

```
ggplot(  
  data = beaches,  
  aes(x = date, y = temperature, group=1)  
) +  
  geom_line(color='red', size=2) + ###  
  geom_point(color='black', size = 3) ###
```

Order of laying down geometries matters

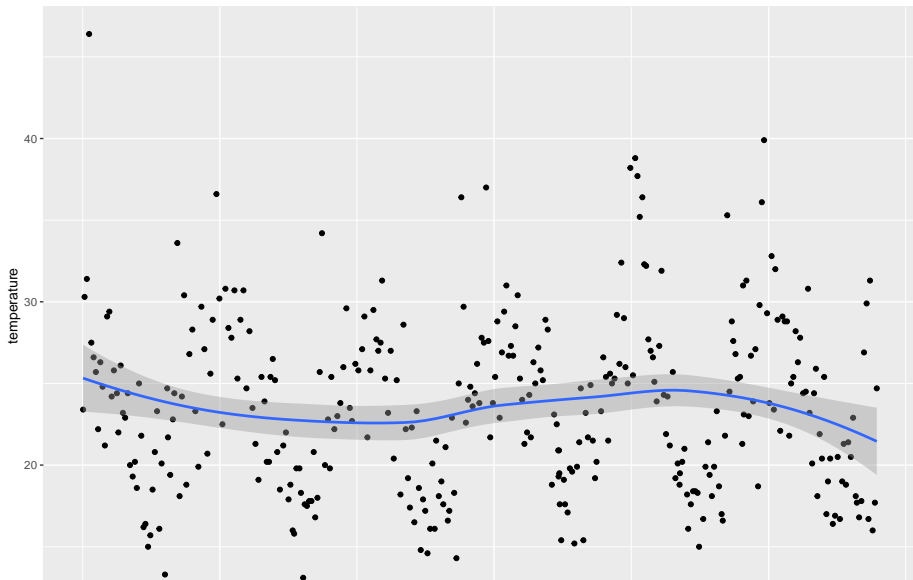


Doing some computations

```
ggplot(  
  data = beaches,  
  aes(x = date, y = temperature, group=1)  
) +  
  geom_point() +  
  geom_smooth(method='loess')
```

Averages over 75% of the data

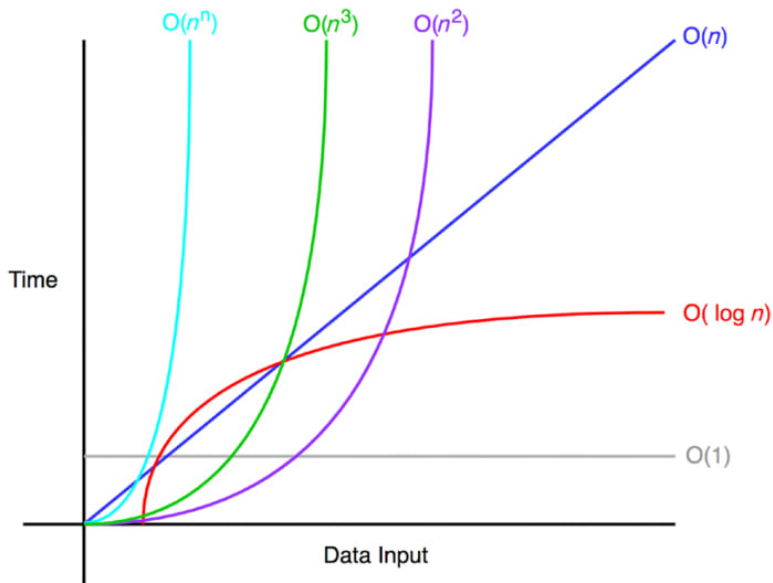
```
#> `geom_smooth()` using formula 'y ~ x'
```



Doing some computations

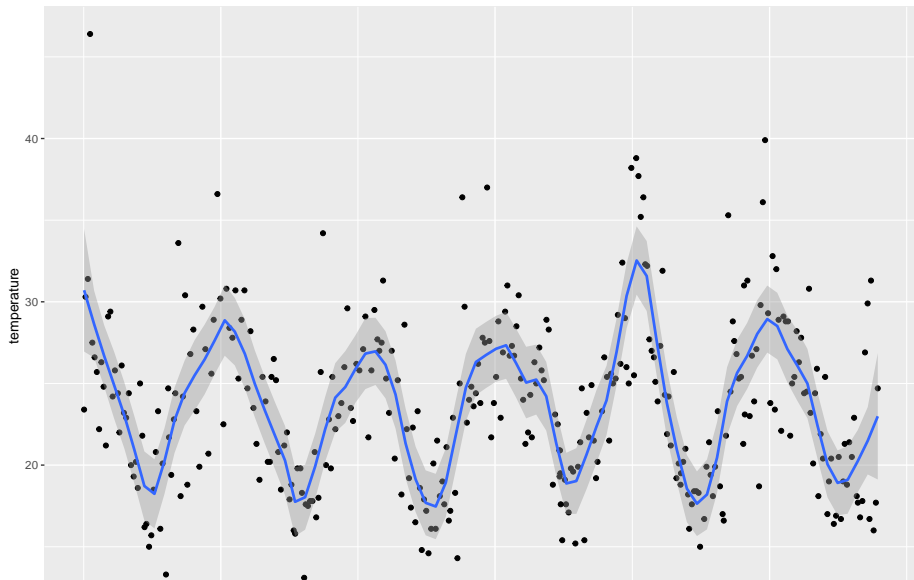
```
ggplot(  
  data = beaches,  
  aes(x = date, y = temperature, group=1)  
) +  
  geom_point() +  
  geom_smooth(method="loess", span = 0.1) ###  
##?geom_smooth, kinda funny...?
```

Big O!



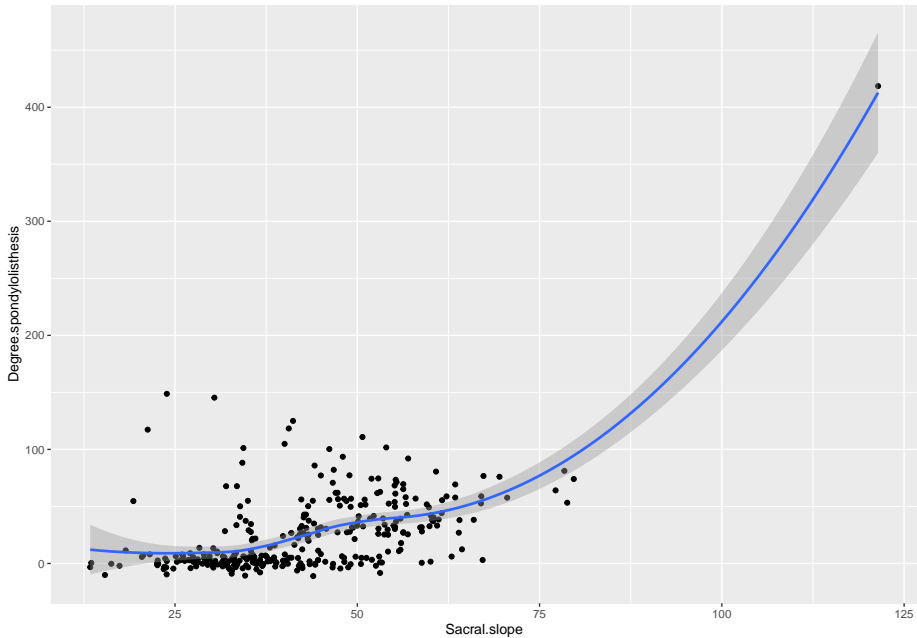
Averages over 10% of the data

```
#> `geom_smooth()` using formula 'y ~ x'
```



Computations over groups

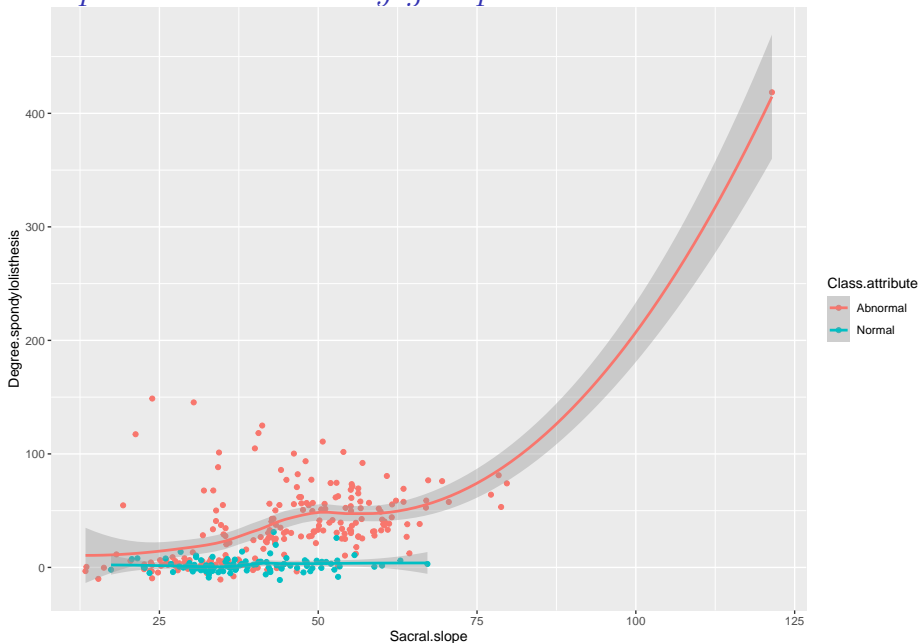
```
ggplot(  
  data = dat_spine,  
  aes(x = Sacral.slope,  
      y = Degree.spondylolisthesis)  
) +  
  geom_point() +  
  geom_smooth()
```



Computations over groups

```
ggplot(  
  data = dat_spine,  
  aes(x = Sacral.slope,  
      y = Degree.spondylolisthesis,  
      color = Class.attribute) ##  
) +  
  geom_point() +  
  geom_smooth()
```

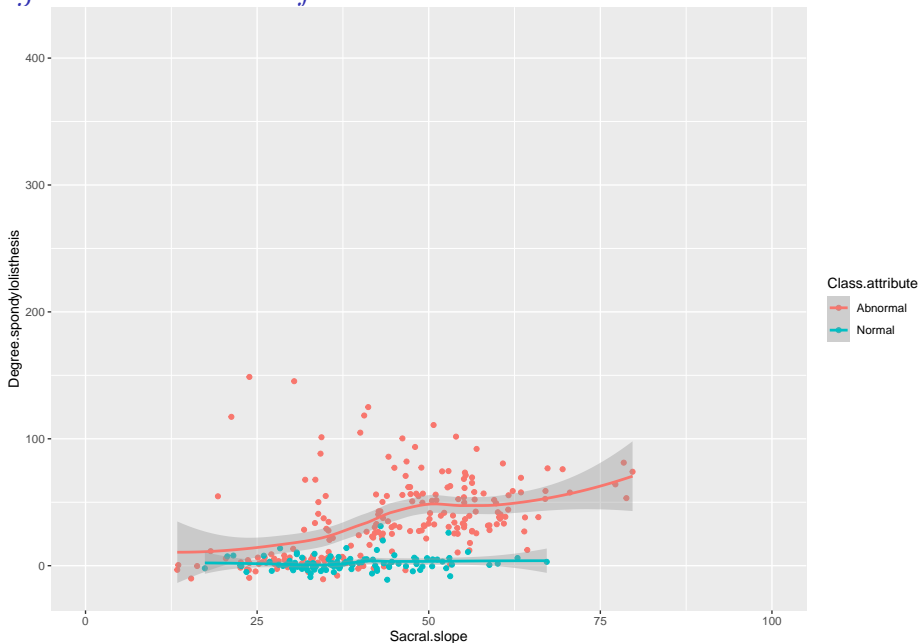
Computation is done by groups



Computations over groups

```
ggplot(  
  data = dat_spine,  
  aes(x = Sacral.slope,  
      y = Degree.spondylolisthesis,  
      color = Class.attribute)  
) +  
  geom_point() +  
  geom_smooth() +  
  xlim(0, 100)#Changing the demonsions of the graphic
```


Ignore the outlier for now

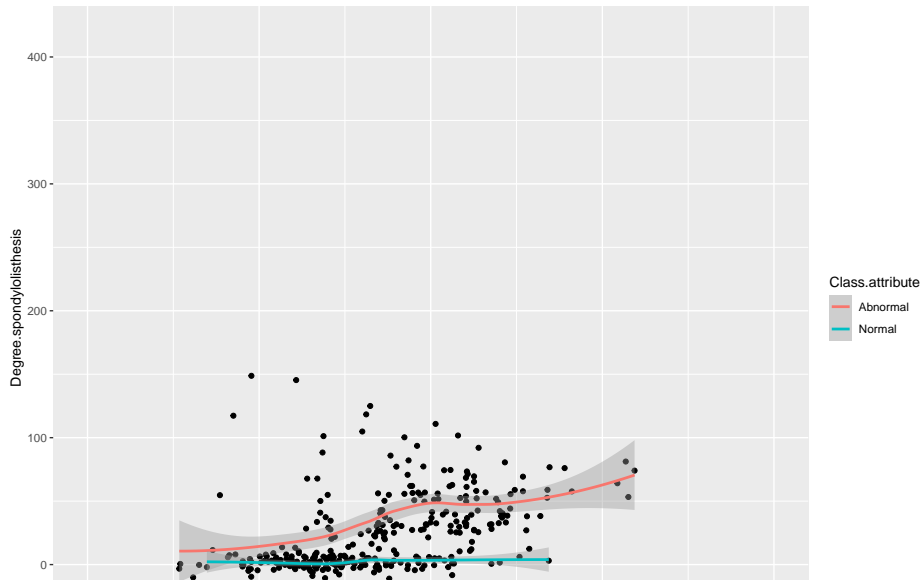


Computations over groups

```
ggplot(  
  data = dat_spine,  
  aes(x = Sacral.slope,  
      y = Degree.spondylolisthesis)  
) +  
  geom_point() +  
  geom_smooth(aes(color = Class.attribute)) + #  
  xlim(0, 100)
```

Only color-code the smoothers

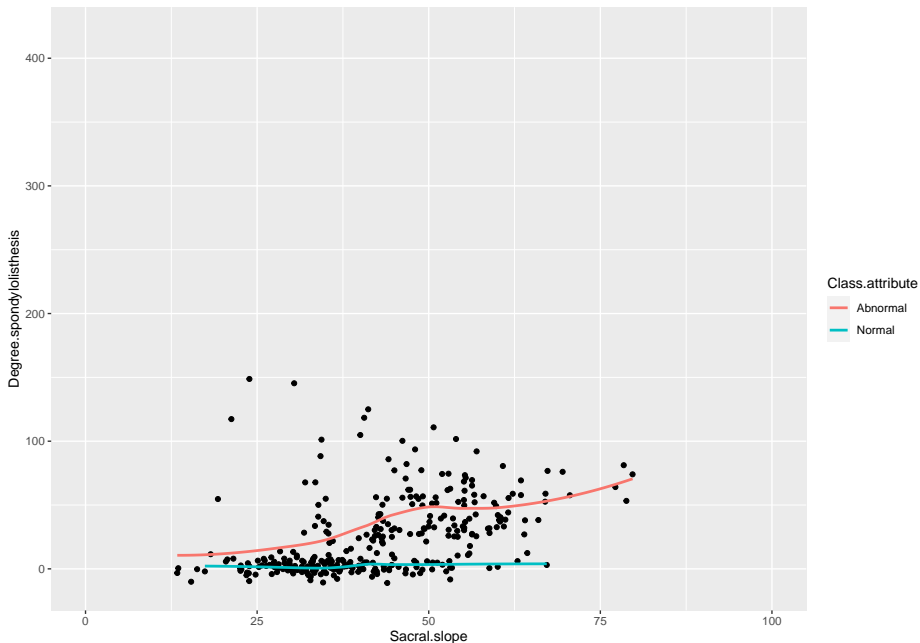
You can change the plot based on where you map the aesthetic



Computations over groups

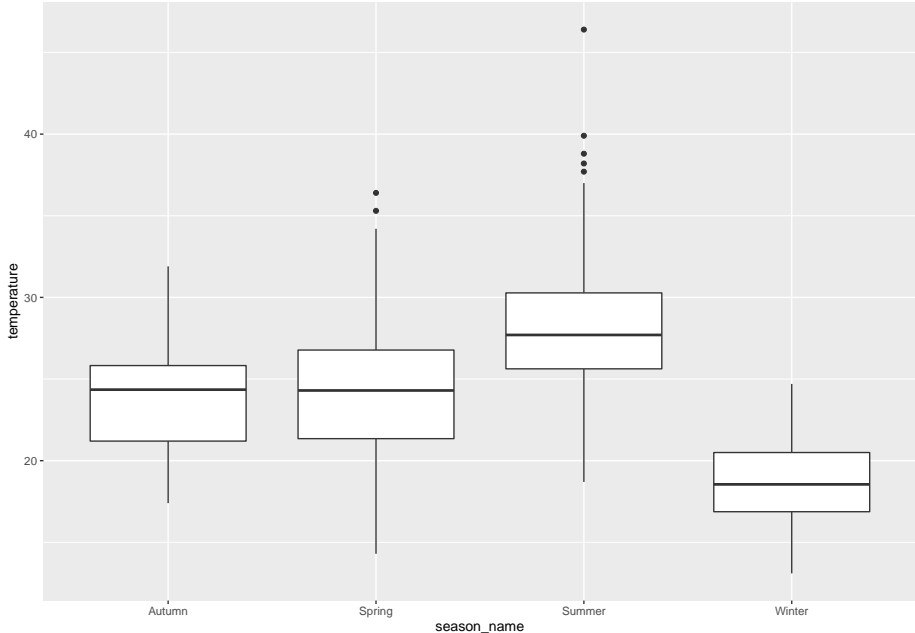
```
ggplot(  
  data = dat_spine,  
  aes(x = Sacral.slope,  
      y = Degree.spondylolisthesis)  
) +  
  geom_point() +  
  geom_smooth(aes(color = Class.attribute),  
              se = F) +  
  #Turning off the confidence interval  
  xlim(0, 100)
```

Looks a little cleaner



Box Plots

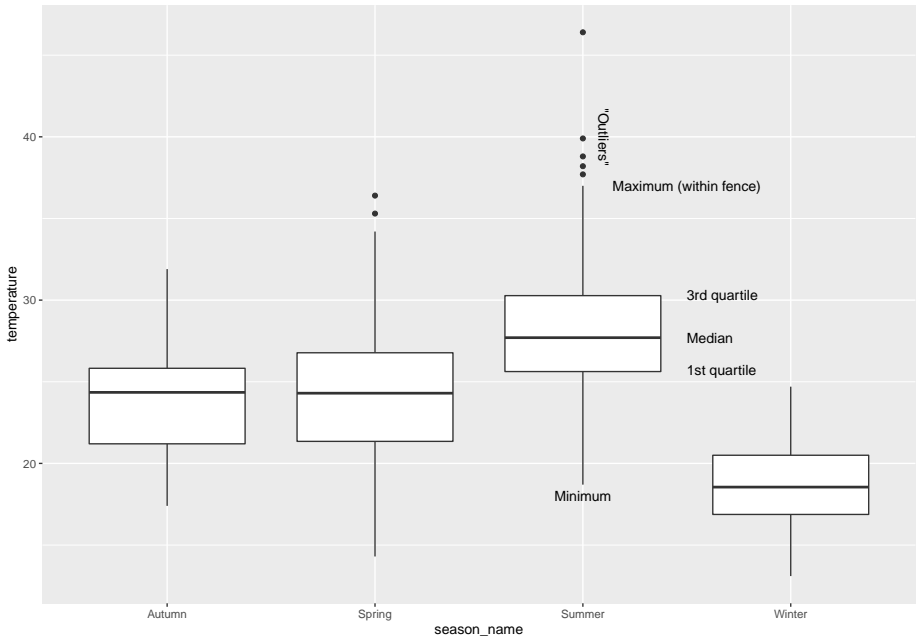
```
ggplot(  
  data = beaches,  
  aes(x = season_name,  
       y = temperature)  
) +  
  geom_boxplot()
```



are the components of a boxplot?

Box Plots

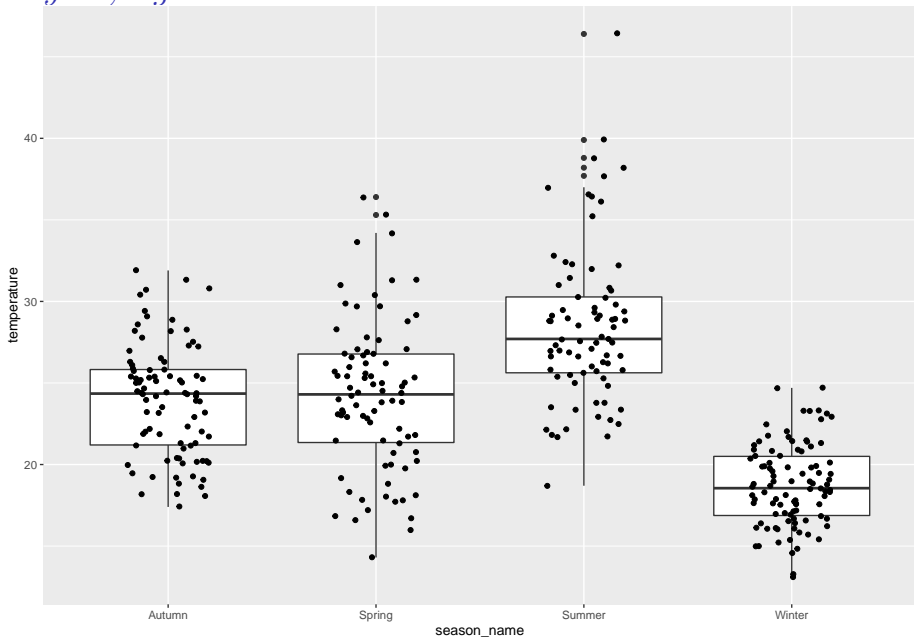
```
ggplot(  
  data = beaches,  
  aes(x = season_name,  
      y = temperature)  
) +  
  geom_boxplot()
```

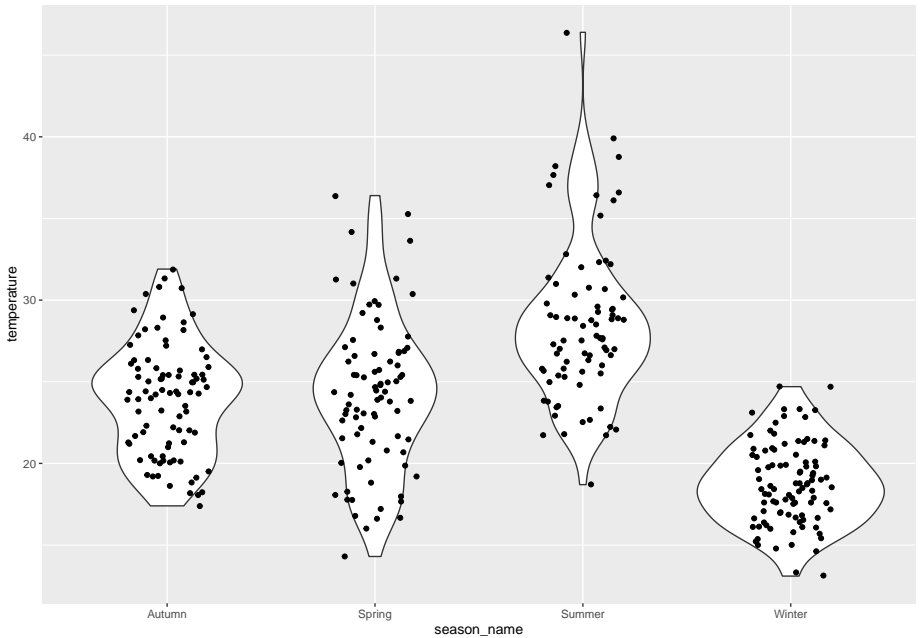



Layers, again

```
ggplot(  
  data = beaches,  
  aes(x = season_name,  
       y = temperature)  
) +  
  geom_boxplot() +  
  geom_jitter(width = 0.2)
```

Layers, again

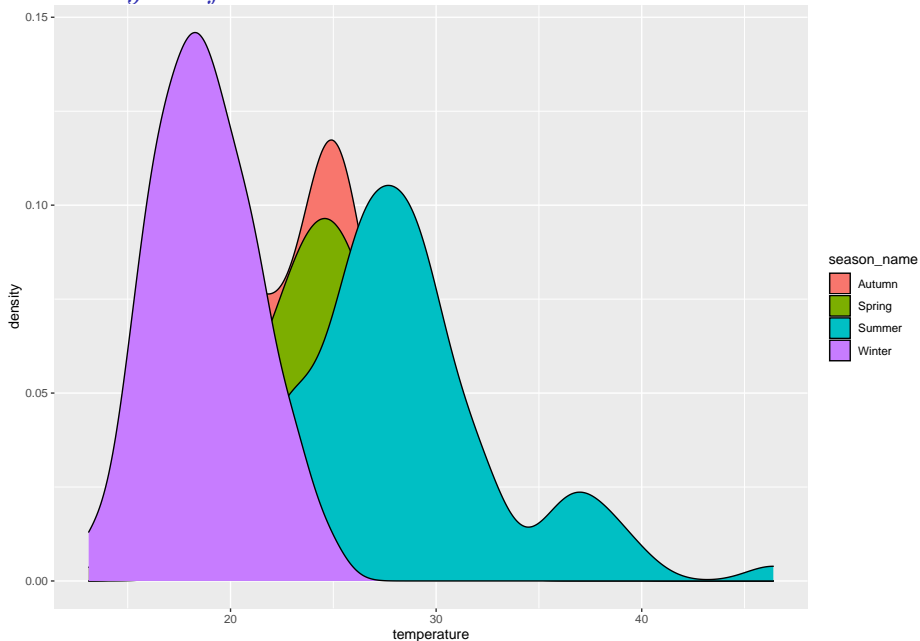




Exploring grouped data

```
ggplot(  
  data = beaches,  
  aes(x = temperature,  
       fill = season_name)  
) +  
  geom_density()
```

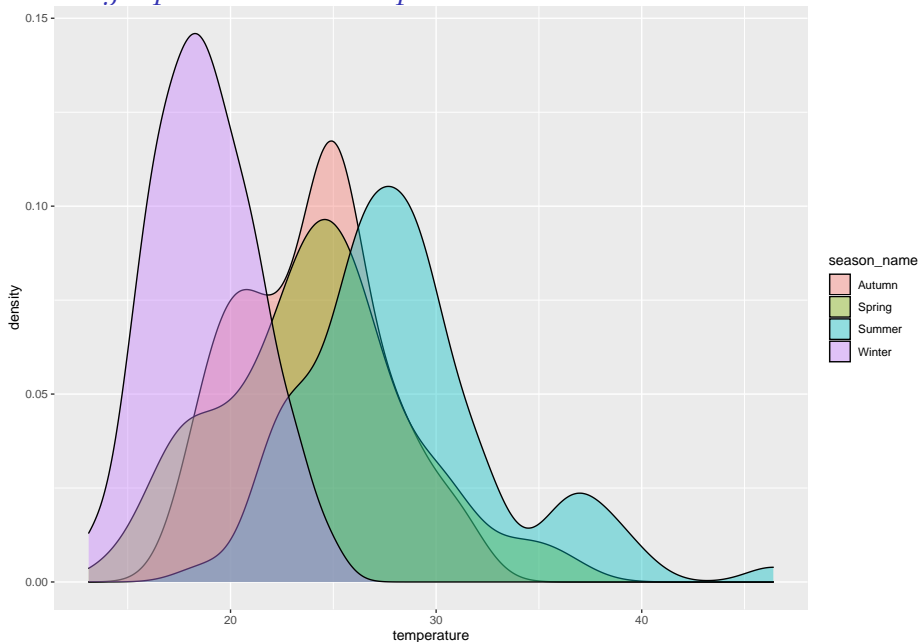
Not very useful



Overlaying graphs

```
ggplot(  
  data = beaches,  
  aes(x = temperature,  
       fill = season_name)  
) +  
  geom_density(alpha = 0.4) # Changes the transparency
```

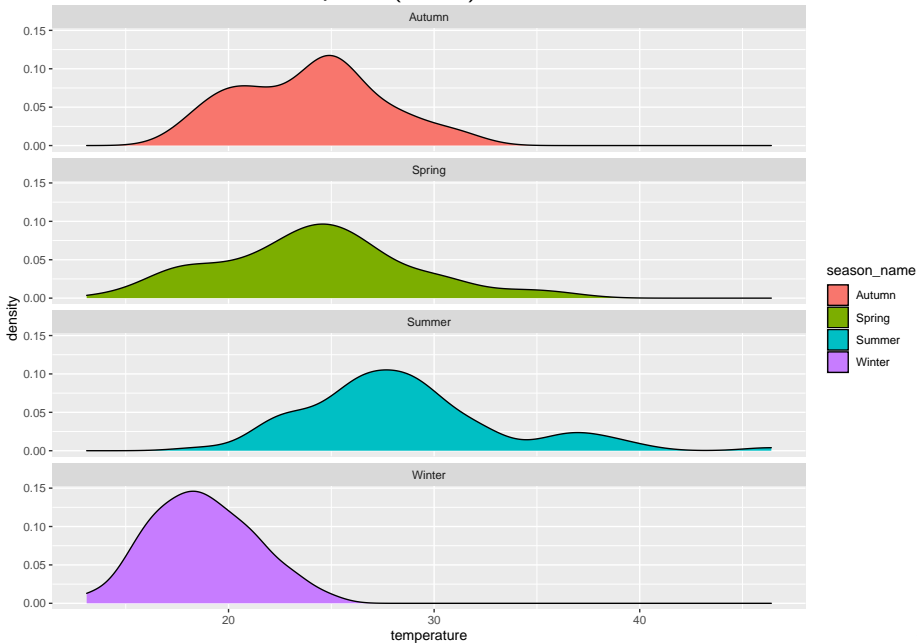
Make graphs more transparent



Exploding graphs

```
ggplot(  
  data = beaches,  
  aes(x = temperature,  
      fill = season_name)  
) +  
  geom_density() +  
  facet_wrap(~ season_name, ncol = 1) ###
```

This is called “small multiples” (Tufte)

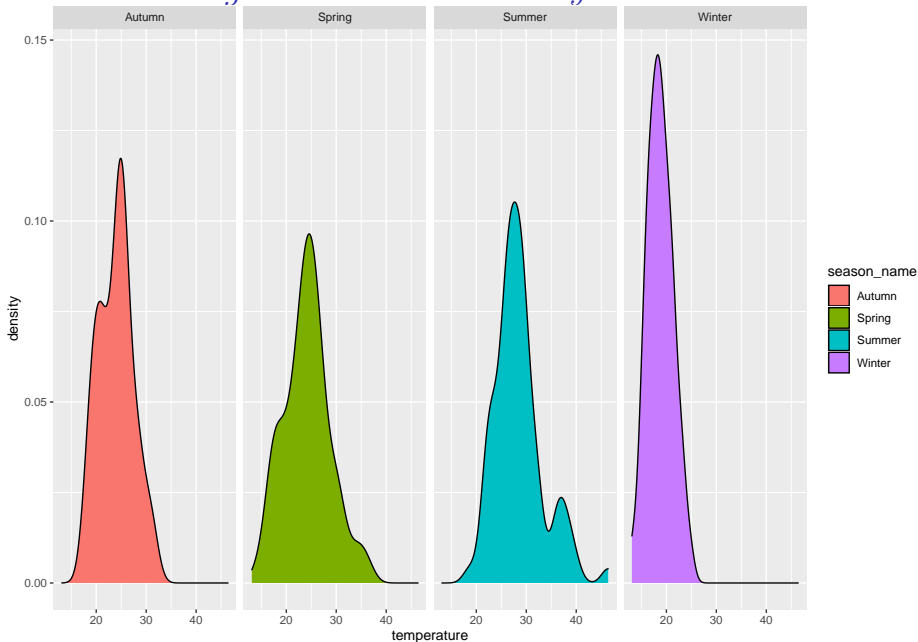


Notice that all the graphs have the same x-axis. This is a good thing

Exploding graphs

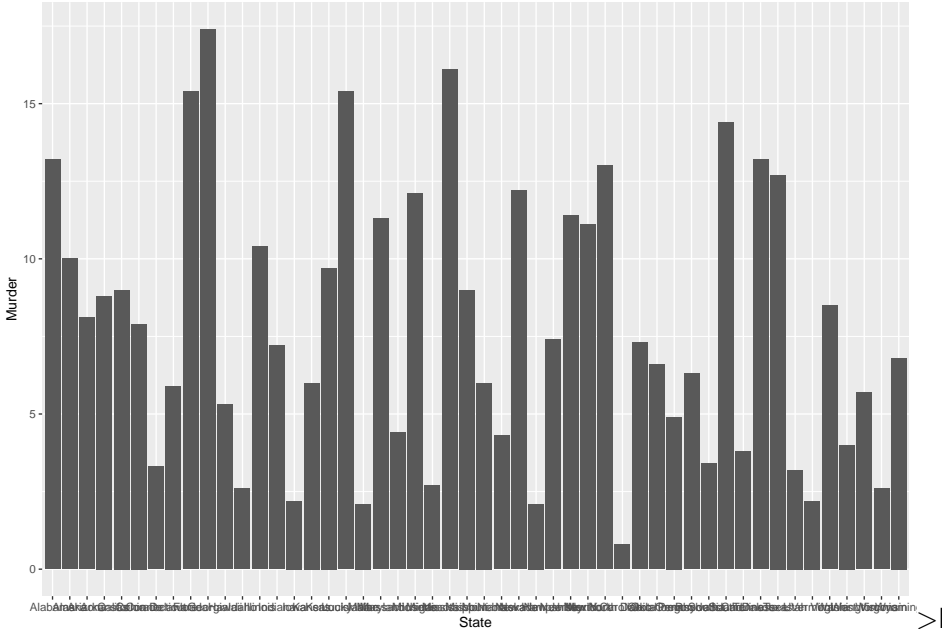
```
ggplot(  
  data = beaches,  
  aes(x = temperature,  
       fill = season_name)  
) +  
  geom_density() +  
  facet_wrap(~ season_name, nrow = 1) ###
```

We can arrange them the other way too



Order and orientation: Arrests in the USA in 1973

```
arrests <- import('USArrests.csv')
ggplot(
  data = arrests,
  aes(x = State,
      y = Murder)
) +
  geom_bar(stat = 'identity')
```

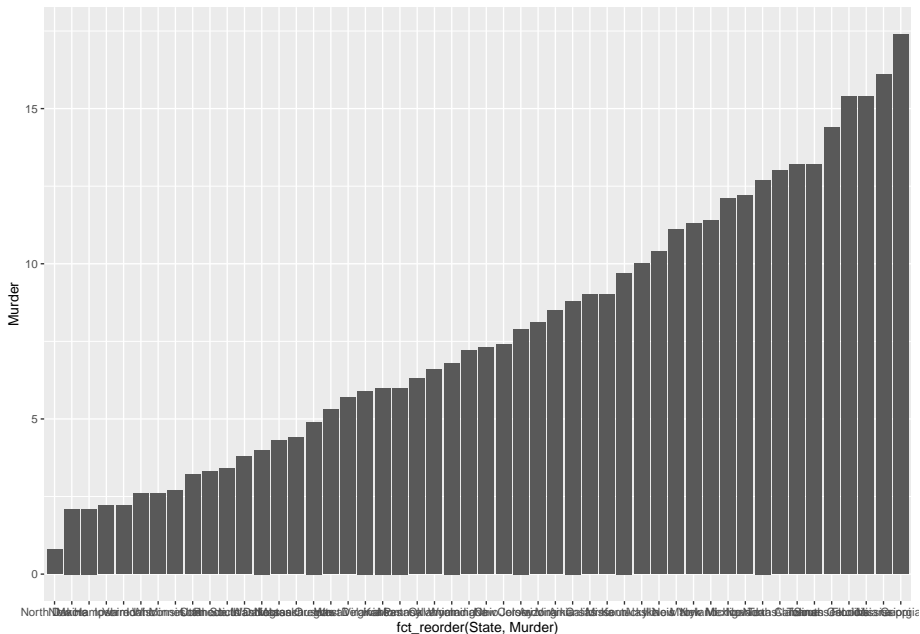


to read, there is no ordering, and x-labels can't be seen

Arrests in the USA in 1973

```
ggplot(  
  data = arrests,  
  aes(x = fct_reorder(State, Murder), #Order by murder rate  
      y = Murder)  
) +  
  geom_bar(stat = 'identity')
```

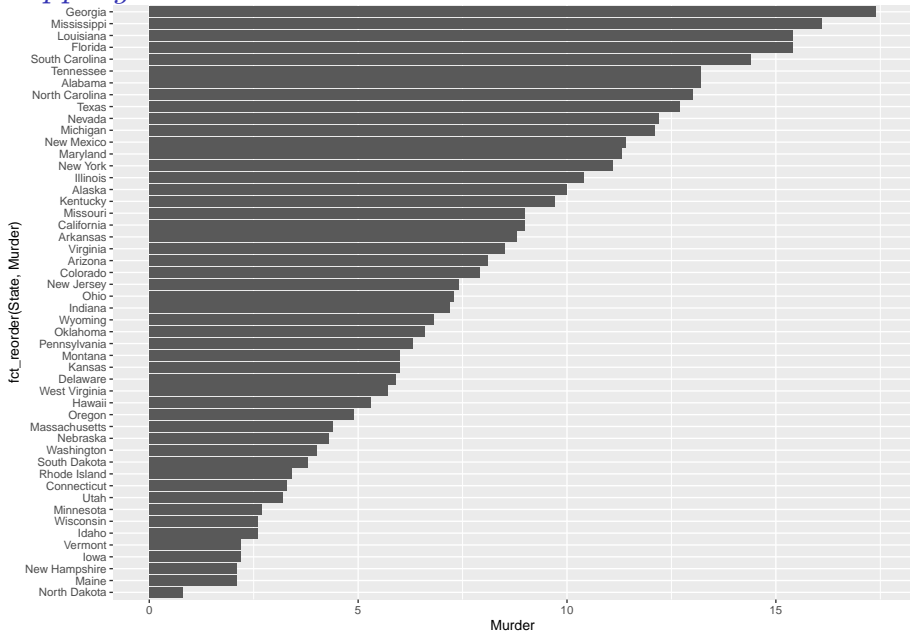
Arrest in the USA in 1973



Arrests in the USA in 1973

```
ggplot(  
  data = arrests,  
  aes(x = fct_reorder(State, Murder), # Order by murder rate  
      y = Murder)  
) +  
  geom_bar(stat = 'identity') +  
  coord_flip() #Flipping the coordinates
```

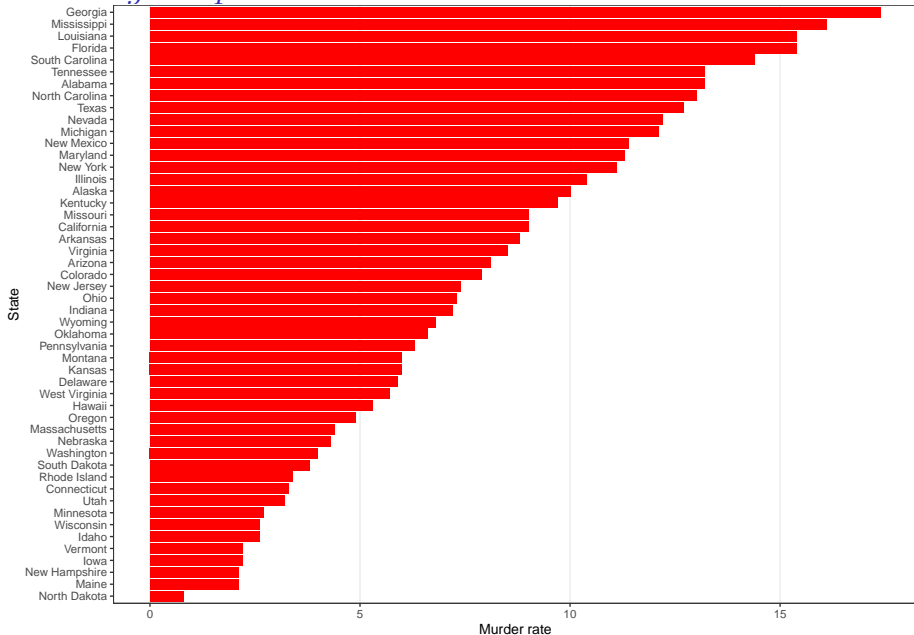
Flipping the axes makes the states readable



Arrests in the USA in 1973

```
ggplot(  
  data = arrests,  
  aes(x = fct_reorder(State, Murder), # Order by murder rate  
      y = Murder)  
) +  
  geom_bar(stat = 'identity', fill="red") +  
  labs(x = 'State', y = 'Murder rate') + # Adding labels  
  theme_bw() + # Theme  
  theme(panel.grid.major.y = element_blank(), #  
        panel.grid.minor.x = element_blank()) +  
  coord_flip() # Flip last
```

Cleaning it up a little



Themes

ggplot comes with a default color scheme. There are several other schemes available

- `scale*_brewer` uses the ColorBrewer palettes
- `scale*_gradient` uses gradients
- `scale*_distill` uses the ColorBrewer palettes, for continuous outcomes

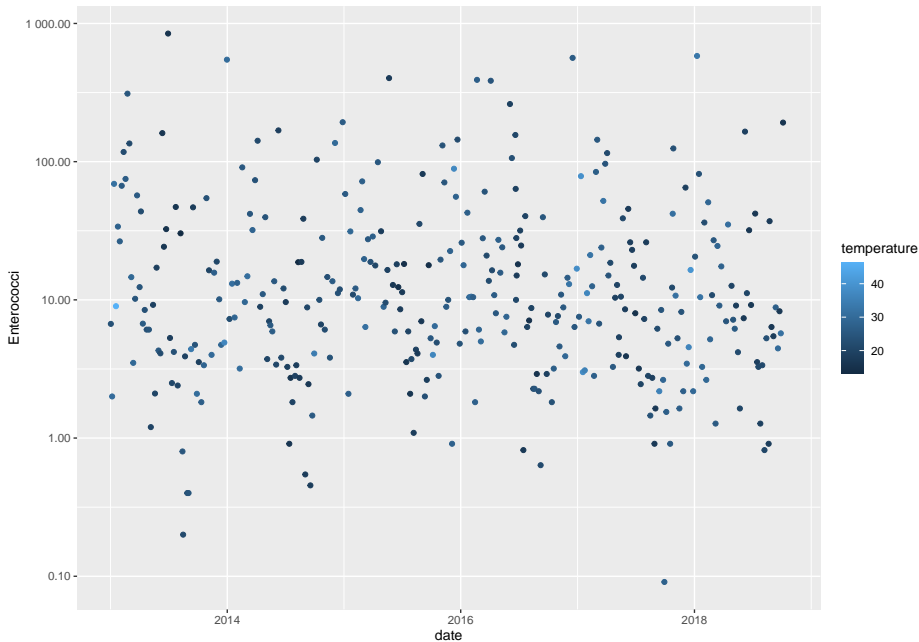
*Here * can be color or fill, depending on what you want to color*

Note color refers to the outline, and fill refers to the inside

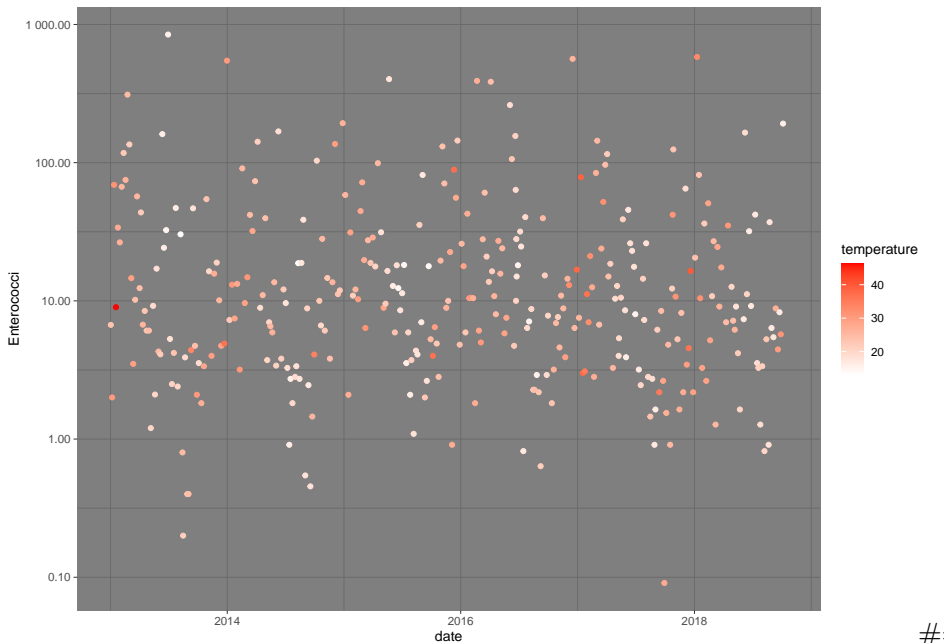
No Theme

```
ggplot(  
  data = beaches,  
  aes(x = date, y = enterococci,  
      color = temperature)  
) +  
  geom_point() +  
  scale_y_log10(name = 'Enterococci',  
               label = scales::number_format(digits=3))
```

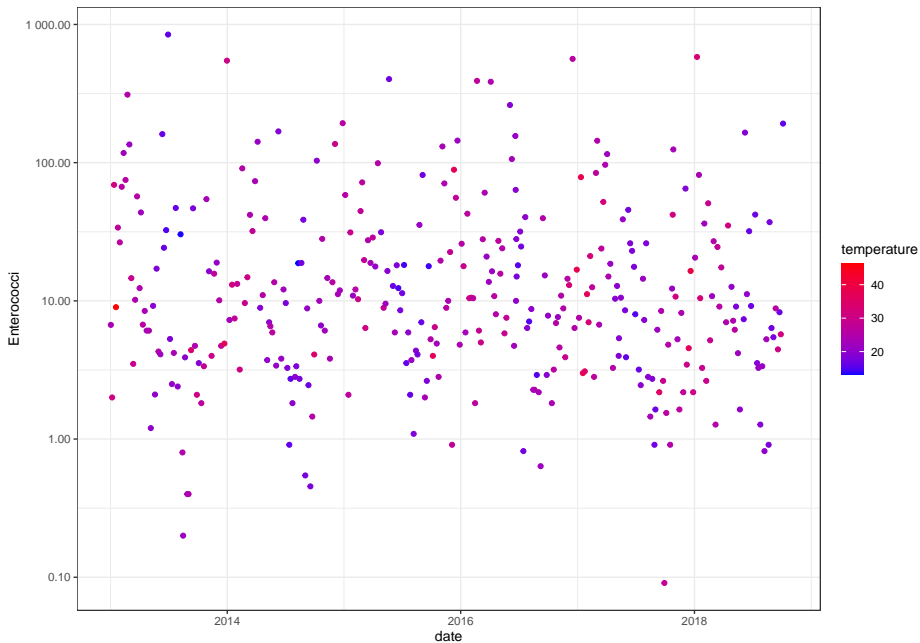
No Theme



Dark



Black and White



Themes

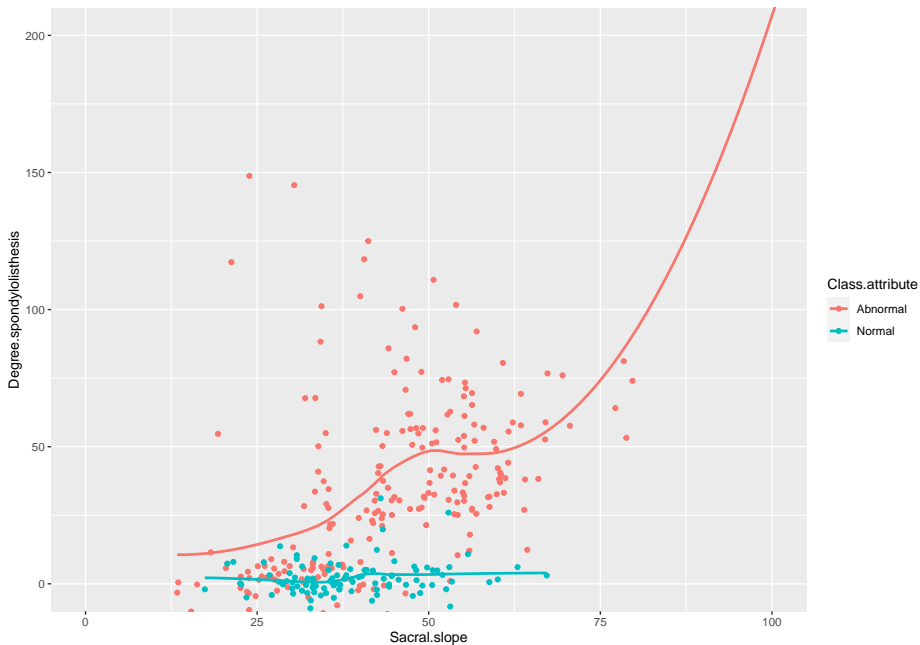
You can create your own custom themes to keep a unified look to your graphs

ggplot comes with

- `theme_classic`
- `theme_bw`
- `theme_void`
- `theme_dark`
- `theme_gray`
- `theme_light`
- `theme_minimal`

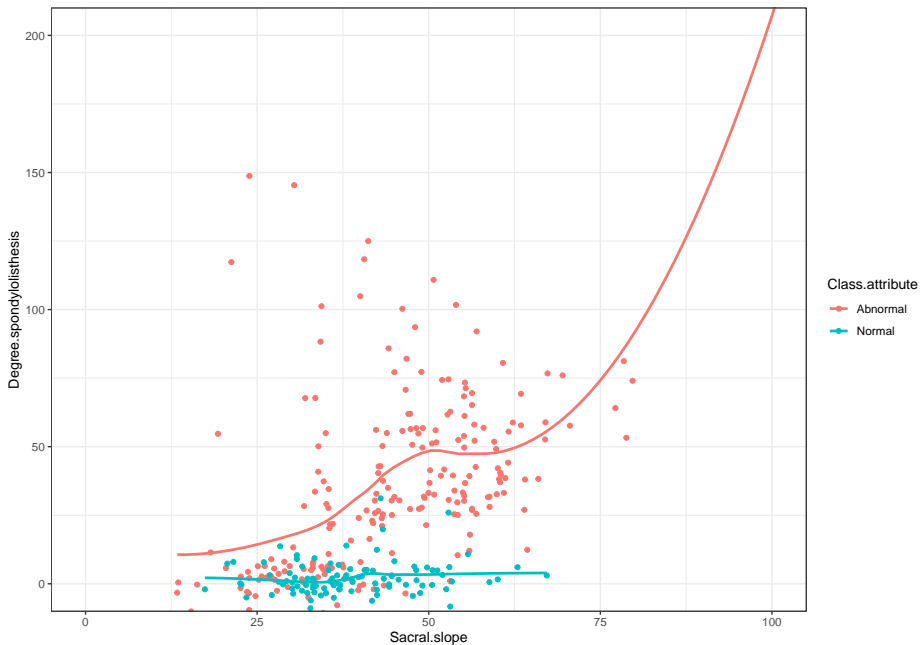
Create your own

```
ggplot(  
  data = dat_spine,  
  aes(x = Sacral.slope, y = Degree.spondylolisthesis,  
      color = Class.attribute)  
) +  
  geom_point() +  
  geom_smooth(se = F) +  
  coord_cartesian(xlim = c(0, 100),  
                  ylim = c(0, 200))
```



Create your own

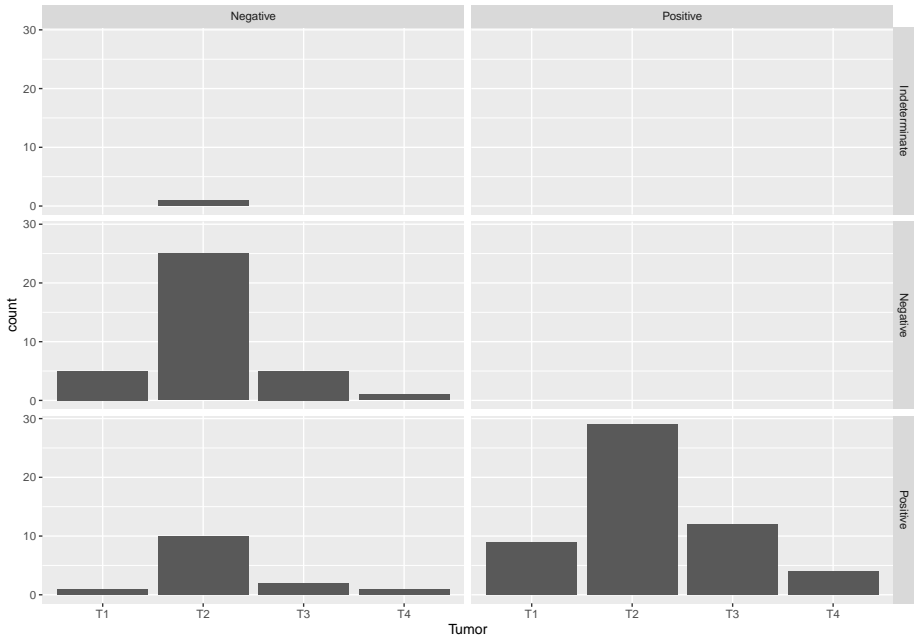
```
my_theme <- function(){  
  theme_bw()  
}  
  
ggplot(  
  data = dat_spine,  
  aes(x = Sacral.slope, y = Degree.spondylolisthesis,  
      color = Class.attribute)  
) +  
  geom_point() +  
  geom_smooth(se = F) +  
  coord_cartesian(xlim = c(0, 100),  
                  ylim = c(0,200)) +  
  my_theme() # Just Black and White
```



Create your own

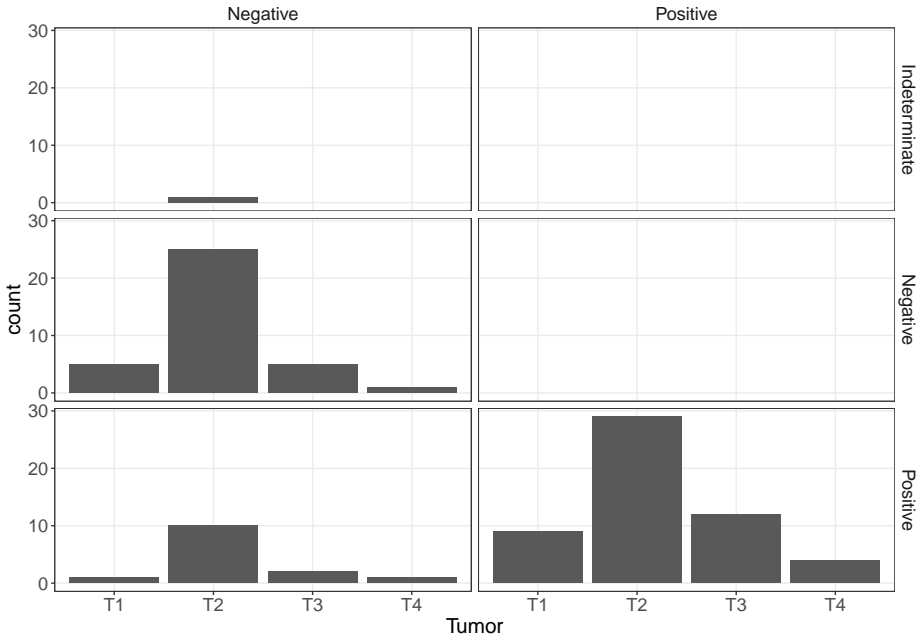
```
my_theme <- function(){  
  theme_bw() +  
  theme(axis.text = element_text(size = 14),  
        axis.title = element_text(size = 16),  
        panel.grid.minor = element_blank(),  
        strip.text = element_text(size=14),  
        strip.background = element_blank())  
}
```

```
ggplot(  
  data = dat_brca,  
  aes(x = Tumor)) +  
  geom_bar() +  
  facet_grid(rows = vars(ER.Status),  
             cols = vars(PR.Status))
```



Create your own

```
ggplot(  
  data = dat_brca,  
  aes(x = Tumor)  
)+  
geom_bar() +  
facet_grid(rows = vars(ER.Status),  
            cols = vars(PR.Status)) +  
my_theme()
```



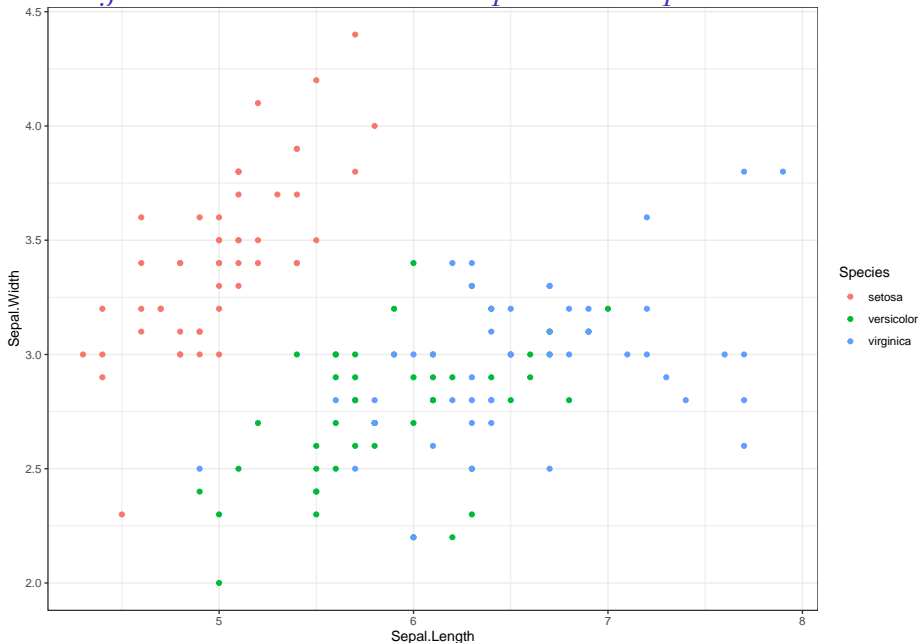
Annotations: Stand-alone stories

- Data visualization to stand on its own
- Relevant information should be placed on the graph
- However, you need to balance the information content with real estate
 - ▶ Don't clutter the graph and make it not readable

Adding derived statistics to a plot: Group means

```
ggplot(iris,  
  aes(x = Sepal.Length,  
    y = Sepal.Width,  
    color = Species)  
)+  
geom_point()+  
theme_bw()
```

Adding derived statistics to a plot: Group means



Adding derived statistics to a plot: Group means

```
means <- iris %>% group_by(Species) %>%  
  summarize_at(vars(starts_with('Sepal')),  
               mean)
```

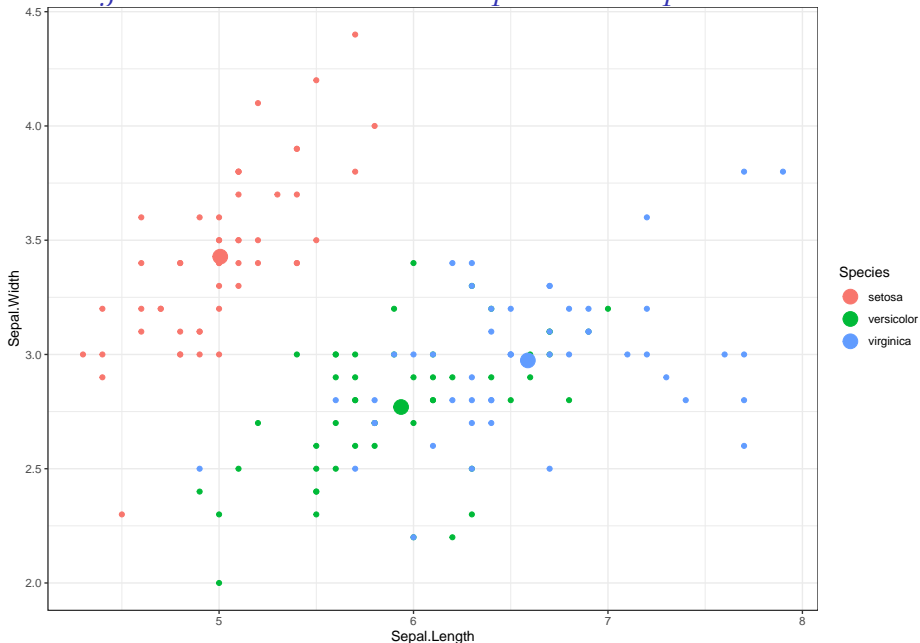
```
means
```

```
#> # A tibble: 3 x 3  
#>   Species      Sepal.Length Sepal.Width  
#>   <fct>          <dbl>          <dbl>  
#> 1 setosa         5.01            3.43  
#> 2 versicolor    5.94            2.77  
#> 3 virginica      6.59            2.97
```

Adding derived statistics to a plot: Group means

```
ggplot(iris,  
      aes(x = Sepal.Length,  
          y = Sepal.Width,  
          color = Species))  
  ) +  
  geom_point() +  
  geom_point(data = means,  
            size=5) +  
  theme_bw()
```

Adding derived statistics to a plot: Group means



Adding regression metrics

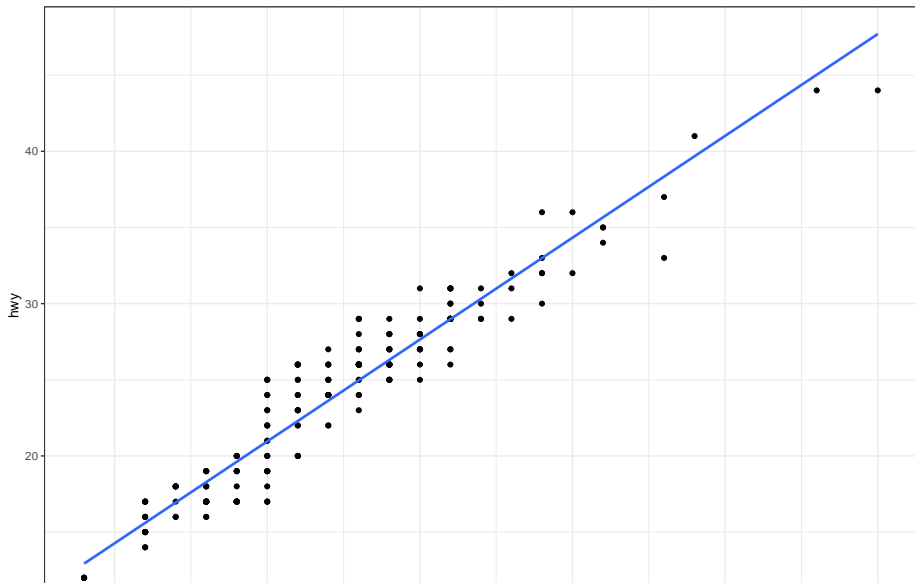
Regress highway mileage on city mileage (data: mpg)

```
mod1 <- lm(hwy ~ cty, data = mpg)
r2 <- broom::glance(mod1) %>% pull(r.squared)

ggplot(mpg,
       aes(x = cty, y = hwy)
) +
  geom_point() +
  geom_smooth(method = 'lm', se=F) +
  theme_bw()
```

Adding regression metrics

```
#> `geom_smooth()` using formula 'y ~ x'
```



Regress highway mileage on city mileage

```
mod1 <- lm(hwy ~ cty, data = mpg)
r2 <- broom::glance(mod1) %>% pull(r.squared) %>% #pull is part of
  round(., 2) #part of base R, rounding behind the . by 2

ggplot(mpg,
       aes(x = cty, y = hwy)) +
  geom_point() +
  geom_smooth(method = 'lm', se=F) +
  annotate(geom='text',
         x = 15, y = 40,
         label=glue::glue("R^2 == {r}",r=r2),
         size=6,
         parse=T) +
  theme_bw()
```

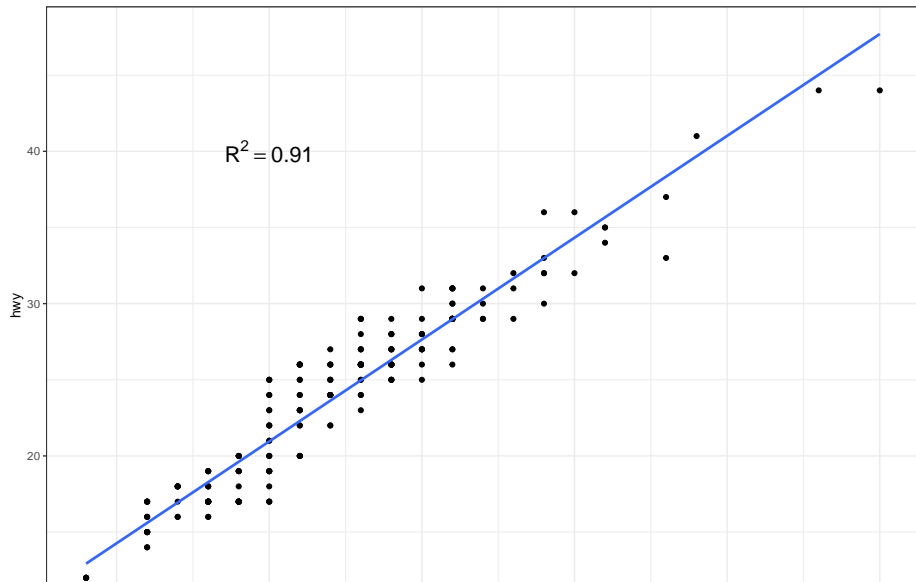
Glance creates a quick summary of the model

```
broom::glance(mod1)
```

```
#> # A tibble: 1 x 12
#>   r.squared adj.r.squared sigma statistic  p.value    df logLik   AIC   BIC
#>   <dbl>         <dbl> <dbl>     <dbl>    <dbl> <dbl> <dbl> <dbl> <dbl>
#> 1     0.914         0.913  1.75     2459. 1.87e-125     1 -462.  931.  941.
#> # ... with 3 more variables: deviance <dbl>, df.residual <int>, nobs <int>
```

Nice Addition

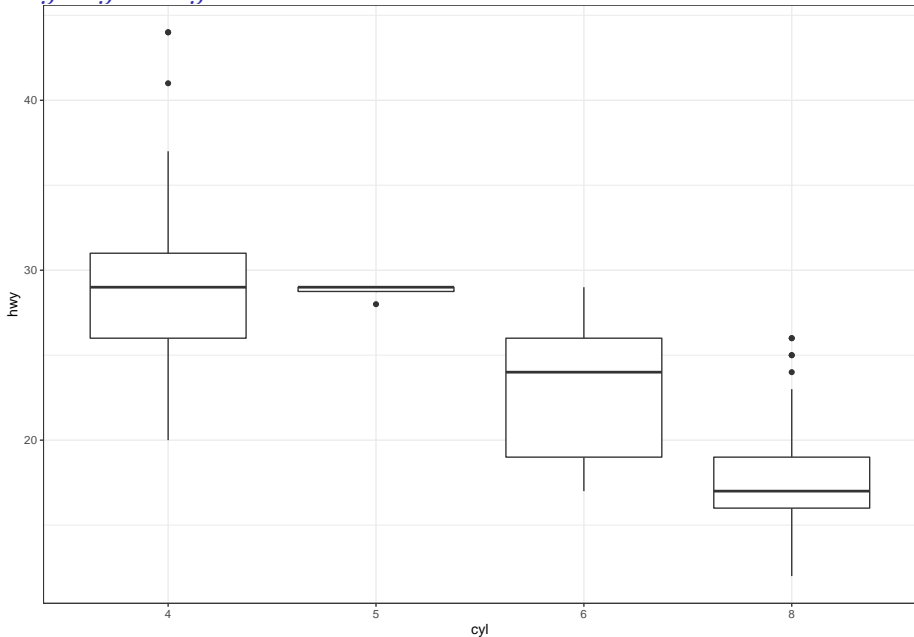
```
#> `geom_smooth()` using formula 'y ~ x'
```



Highlighting regions

```
mpg %>%  
  mutate(cyl = as.factor(cyl)) %>%  
  ggplot(aes(x = cyl, y = hwy)) +  
  geom_boxplot() +  
  theme_bw()
```

Highlight regions



Highlighting regions

```
mpg %>%  
  mutate(cyl = as.factor(cyl)) %>%  
  ggplot(aes(x = cyl, y = hwy))  
    +  
  geom_boxplot() +  
  theme_bw() +  
  annotate(geom = 'rect',  
    xmin=3.75,xmax=4.25,  
    ymin = 22, ymax = 28,  
    fill = 'red',  
    alpha = 0.2) +  
  annotate('text',  
    x = 4.5, y = 25,  
    label = 'Outliers?',  
    hjust = 0) +  
  coord_cartesian(xlim = c(0,5)) +  
  theme_bw()
```


Highlighting regions

